

DOCUMENT 33 · AI RISK METHODOLOGY

AI risk methodology

How AI risks are identified, assessed, treated, monitored and reported in each initiative and in the portfolio

Fernando García Varela Version 0.1 (working draft) 16-09-2026 Draft for review

10

risk categories

70

typical risks
catalogued

5 × 5

likelihood and
impact matrix

5

impact axes

IN THIS DOCUMENT

01	Purpose and scope	08	Risk register
02	Principles	09	Catalogue of typical risks
03	The process	10	Portfolio risk
04	Assessment scales	11	Risk appetite and key risk indicators
05	Inherent risk, controls and residual risk	12	Risk reporting
06	Risk responses	13	Associated tools and templates
07	Acceptance and escalation	14	Related documents
		15	Version control

Legal notice and disclaimer. SEVEN-G is a reference methodological framework provided "as is" and for information purposes only. It does not constitute legal, regulatory, financial or professional advice, nor does it guarantee compliance with any law or standard. References to general regulation (such as the EU AI Act, the GDPR, DORA or NIS2), to technical standards and to regulation specific to each sector or jurisdiction may be incomplete, may not apply to a particular case or may become out of date as a result of regulatory changes, interpretations or supervisory positions after their consultation date. **Each organisation that uses SEVEN-G is solely responsible for identifying the regulation that applies to it, verifying that it is current and certifying its own regulatory compliance**, with the appropriate qualified advice. The author accepts no liability for any use made of this content or for any decisions taken on the basis of it. The data, figures, companies and cases in the examples are fictitious or illustrative.

THE ESSENTIALS

Always mandatory

Risk register for each initiative from phase 3, with the 5 × 5 scale, inherent and residual risk, and acceptance by the body that corresponds to its level. A Critical residual risk without board approval blocks G3 and G5. Portfolio risk and key risk indicators are applied in proportion to the size of the portfolio.

Obligation matrix and layered reading (document 94)

01 Purpose and scope

This document defines the methodology with which an organisation that applies SEVEN-G manages the risks of its AI systems and initiatives. It is the reference for the framework's risk scales: documents

13 (risk appetite), 14 (portfolio), 21 (*gate* criteria), 35 (security), 36 (third parties), templates P12 and P13 and tool T06 use exactly the scales and levels set here.

1.1 What it covers

- The six-step risk management process: identify, analyse, evaluate, treat, monitor and report.
- Its application at the framework's two levels: the **initiative** (phases 0–7) and the **portfolio** (corporate cycle C1–C5).
- The likelihood and impact scales, the 5 × 5 matrix and the risk levels.
- Inherent and residual risk, and the assessment of control effectiveness.
- Risk acceptance, escalation and the risk register.
- The **catalogue of typical risks** `RT-<CAT>-NN`, organised into ten categories.
- Portfolio risk: concentration, correlation and supplier dependency.
- Key risk indicators.

1.2 What it does not cover

- The company's appetite and tolerances, which are approved in C2 (document 13).
- The detail of security and agent controls (document 35), third-party controls (document 36) and regulatory compliance controls (document 34).
- The management of incidents and nonconformities (document 37). A materialised risk is managed as an incident; the risk methodology captures its lessons.

1.3 Reference standards

The methodology is consistent with [ISO/IEC 23894](#) (guidance on AI risk management, which adapts the [ISO 31000](#) process), with the risk assessment and treatment requirements of [ISO/IEC 42001](#), with the *Map*, *Measure* and *Manage* functions of the [NIST AI RMF](#) and its generative AI profile [NIST AI 600-1](#), and with the risk management system that [Regulation \(EU\) 2024/1689](#) requires of providers of high-risk systems. SEVEN-G does not replace the company's corporate risk framework: it specialises it for AI and must be integrated into it (same level language, same bodies, same report).

This document does not constitute legal advice.


Principles

#	Principle	Practical consequence
1	Risk is assessed before investing	The full assessment is mandatory evidence for G3, the main stop gate.
2	An untested control does not reduce risk	Residual risk is only considered verified when the effectiveness of the control is evidenced (normally in phase 5).
3	The worst impact axis is assessed	Impact is the highest of the five axes; it is not averaged.
4	Every risk has an owner	Without a named owner, the risk is not registered.
5	Those who build do not accept their own risk	Acceptance follows the scale in section 7 and respects segregation of duties (01 §8).
6	Risk is systemic	In addition to the risk of each initiative, portfolio concentration and correlation are managed (principle 5 of 01 §3).
7	The register is living	It is updated at every <i>gate</i> , at every continuity review, after every incident and whenever there is a relevant change.

The process

AI risk management process

Six steps repeated at every gate and every review



Source: SEVEN-G

Step	What is done	Output	Owner
1. Identify	Go through the catalogue of typical risks (section 9), the phase 0 context, the regulatory classification, the architecture and the suppliers. Describe each risk as cause → event → consequence .	List of risks with the associated RT code.	AI Product Owner with the AI Technical Owner; facilitated by the AI Risk Owner.
2. Analyse	Estimate inherent likelihood and impact using the scales in section 4, justifying the determining impact axis. Identify existing controls.	Inherent level per risk.	Initiative team.
3. Evaluate	Compare with the appetite and tolerances (document 13). Decide which risks require treatment and in what order.	Prioritisation and proposed response.	AI Risk Owner.
4. Treat	Choose the response (avoid, mitigate, transfer, accept), design and apply controls, estimate the target residual and obtain acceptance of the residual.	Mitigation and contingency plan (P13); acceptance decision.	Risk owner; acceptance by whoever is competent under section 7.
5. Monitor	Verify the effectiveness of controls, monitor key risk indicators, review at the frequency set for the level and when triggers occur.	Updated register; verified residual.	AI Operations Owner and risk owner; verified by the AI Auditor at the <i>gates</i> .
6. Report	Present the risk profile to the body that decides the <i>gate</i> and, in aggregate, to the AI Committee and the board.	Risk information in P29, board dashboard (T17).	AI Office with the AI Risk Owner.

3.1 In the initiative lifecycle

Phase or gate	Risk activity	Evidence
0 • G0	Constraints and obvious risks (prohibited practice, protected data, critical function). Intensity determination.	Context statement (P02), intensity (P04).
1 • G1	Preliminary identification of risks that could rule out the opportunity.	Screening notes (P06).
2 • G2	Risks of the value hypothesis (RT-ECO, RT-EST) and stop criteria linked to risk.	Hypothesis canvas (P08).
3 • G3	Full assessment: inherent risk, planned controls, target residual, acceptance. No Critical risk without accepted treatment.	Matrix and register (P12); mitigation and contingency plan (P13); supplier assessment (P14).
4 • G4	Every Medium or higher risk has a designed control, with a code where one exists (SEG, AG from document 35).	Security design (P18), oversight design (P17), updated register.

Phase or gate	Risk activity	Evidence
5 • G5	Control effectiveness tested; residual verified ; acceptance ratified.	Validation results (P22), updated register, sign-off (P23).
6 • R6	Monitoring, indicators, incidents, changes; review of the classification.	Updated register; incident register (P27).
7 • G7	Accumulated risks and retirement risks (data, models, dependencies).	Scale or retire decision (P30).

3.2 In the portfolio

Stage	Risk activity
C1 • Diagnosis	Risk profile of the inventoried systems, including those that predate the framework and corporate use of AI.
C2 • Direction	Approval of the appetite, tolerances, economic materiality threshold and indicator thresholds (document 13).
C3 • Portfolio	Concentration, correlation and supplier dependency when prioritising (section 10).
C4 • Oversight	Monthly report to the AI Committee and quarterly report to the board or board committee (section 12).
C5 • Review	Effectiveness of the risk management system, materialised risks, lessons and recalibration of scales.

04

Assessment scales

4.1 Likelihood

Likelihood is estimated for the initiative's horizon or, for systems in production, for the following twelve months. The frequency reference is used for recurring events and the percentage reference for one-off events.

Likelihood	Name	Indicative reference
1	Rare	Less than once in 5 years or < 5 % over the initiative's horizon

Likelihood	Name	Indicative reference
2	Unlikely	Once every 2–5 years or 5–20 %
3	Possible	Once every 1–2 years or 20–50 %
4	Likely	Several times a year or 50–80 %
5	Almost certain	Monthly or more often, or > 80 %

Supporting criteria: precedents in the company or the sector, results of tests and pilots, exposure (users, volume, external access), existence of a motivated attacker, maturity of the technology. For security threats with an active attacker (RT-GEN, RT-SEG) and external exposure, likelihood **should not** be rated below 3 without evidence from adversarial testing.

4.2 Impact across five axes

Impact is assessed on five axes and the **highest** is taken. The determining axis and its justification are recorded.

M is the **economic materiality threshold** that the company approves in C2 (document 13) in proportion to its size. The percentages of M are indicative and are calibrated in C2.

Impact	Economic	People and rights	Regulatory	Operational	Reputational
1 • Negligible	Less than 1 % of M.	No appreciable effect on people.	No non-compliance.	Brief interruption with no effect on customers or relevant processes.	No external repercussions.
2 • Minor	From 1 % to 5 % of M.	Inconvenience to a few people, immediately correctable.	Formal non-compliance that can be remedied without notifying the authorities.	Degradation lasting hours in a non-critical process.	Isolated complaints.
3 • Moderate	From 5 % to 20 % of M.	Reversible adverse effect on a limited group (wrongful treatment, delay, incorrect information).	Non-compliance that requires formal correction and may require notification to the supervisor.	Interruption of a relevant process for one working day or sustained degradation.	Local or sector-level coverage; recurring complaints.

Impact	Economic	People and rights	Regulatory	Operational	Reputational
4 · Major	From 20 % to 100 % of M.	Significant harm: wrong decisions on rights or access to services, exposure of personal data with risk to individuals.	Mandatory notification, formal request or likely sanctioning proceedings.	Interruption of a critical function within its tolerance, or of a relevant process for several days.	National coverage; measurable loss of customers or partners.
5 · Critical	More than M.	Serious or irreversible harm to health, systematic discrimination or infringement of fundamental rights.	Prohibited practice, possible serious incident under the AI Act , serious sanction or loss of authorisation.	Interruption of a critical function beyond its impact tolerance.	Lasting damage to trust; attention from the supervisor and the board.

4.3 Risk level and 5 × 5 matrix

Level = Likelihood × Impact, with four levels: **Low** 1–4 · **Medium** 5–9 · **High** 10–15 · **Critical** 16–25.

Likelihood \ Impact	1 · Negligible	2 · Minor	3 · Moderate	4 · Major	5 · Critical
5 · Almost certain	5 · Medium	10 · High	15 · High	20 · Critical	25 · Critical
4 · Likely	4 · Low	8 · Medium	12 · High	16 · Critical	20 · Critical
3 · Possible	3 · Low	6 · Medium	9 · Medium	12 · High	15 · High
2 · Unlikely	2 · Low	4 · Low	6 · Medium	8 · Medium	10 · High
1 · Rare	1 · Low	2 · Low	3 · Low	4 · Low	5 · Medium

Rules of use:

1. Whole numbers are used. When in doubt between two values, the higher one is taken and the doubt is recorded.
2. **Extreme impact rule:** a risk with impact 5 on the people and rights axis or on the regulatory axis is treated, for acceptance purposes, as at least **High**, even if its likelihood is 1. This prevents rare but unacceptable events from being accepted by the AI Product Owner.
3. Practices prohibited by the AI Act are not assessed: they are **avoided** (01 §6.5).

Inherent risk, controls and residual risk

5.1 Definitions

Concept	Definition
Inherent risk	Level of the risk without considering controls specific to the initiative. General corporate controls are only considered if they are tested and documented in the company.
Target residual risk	Level expected after applying the planned controls. It is used at G3 and G4.
Verified residual risk	Level after checking the effectiveness of the controls with evidence. It is the level required at G5 and monitored at R6.

5.2 Control types

Type	Acts on	Examples
Preventive	Likelihood	Least privilege for agents, data validation, contractual no-training clause.
Detective	Impact, by shortening the time to detection	Drift monitoring, alerts on anomalous actions, sample review.
Corrective or containment	Impact	Kill switch, rollback plan, manual fallback process.
Transfer	Economic impact	Insurance, contractual indemnity. It does not transfer regulatory accountability or harm to people.

5.3 Control effectiveness

Effectiveness	Criterion	Maximum reduction allowed
Effective	Designed, implemented and tested with evidence in the last review period, with no relevant exceptions.	Up to 2 levels of likelihood (preventive) or of impact (detective or corrective).
Partially effective	Implemented and tested with exceptions, incomplete coverage or unsupervised manual dependency.	1 level.

Effectiveness	Criterion	Maximum reduction allowed
Ineffective	The test fails or the control does not operate as designed.	None. A corrective action is opened.
Untested	Designed or implemented without an effectiveness test.	None in the verified residual. It may be used in the target residual.

Rules:

- The total reduction cannot take likelihood below 1 or impact below the impact the event would produce with the corrective control already activated.
- Transfer controls only reduce the economic axis; impact continues to be determined by the highest of the other axes.
- In generative AI and agents, the effectiveness of controls against prompt injection **must** be tested through adversarial testing (document 35, section 8); a policy or an instruction to the model is not, on its own, an effective control.
- An ineffective control on a High or Critical residual risk is also a nonconformity (document 37).

06

Risk responses

Response	When	Conditions
Avoid	Unacceptable risk with no alternative: prohibited practice, Critical with no viable treatment, cost of the control higher than the value.	Change the scope, design, data, supplier or autonomy level; or Stop at the <i>gate</i> .
Mitigate	Usual response for Medium, High and Critical.	Controls with an owner, deadline and target residual. Plan in P13.
Transfer	Significant economic impact with a transfer market.	Document what is transferred and what is not; review insurance exclusions and indemnity limits.
Accept	Residual within appetite or disproportionate mitigation cost.	Decision by the competent level (section 7), with a validity period and review conditions.

Every High or Critical risk under treatment **must** have a **contingency plan**: what is done if it materialises, who activates it and with what trigger. In systems with the ability to act, the plan includes activating the kill switch or the rollback.

Acceptance and escalation

7.1 Who accepts residual risk

Residual level	Accepted by	Conditions
Low	AI Product Owner	With a record.
Medium	AI Sponsor	With clearance from the AI Risk Owner.
High	AI Committee	With a treatment plan, indicators and a maximum acceptance validity until the next continuity review.
Critical	Not accepted. Exceptionally, only the board or its board committee, within the appetite approved in C2.	Rationale, time limit, compensating controls and monitoring at every quarterly meeting. A Critical residual without that approval blocks G3 and G5.

Rules:

- No one accepts a risk of an initiative they are building (01 §8.2).
- Acceptance expires when the level, the system, the supplier, the model, the regulatory classification or the autonomy level changes, and in any case on the recorded validity date.
- Accepting a risk does not exempt compliance with regulation: a legal breach is not accepted as a risk; it is corrected.

7.2 Escalation triggers and extraordinary review

Trigger	Action	Indicative time limit
New Critical residual risk	Notification to the AI Committee; assessment of stopping or suspension.	Same working day
New High residual risk or a one-level increase	Notification to the AI Sponsor and the AI Risk Owner; review at the next committee meeting.	5 working days
Key risk indicator outside threshold	Review of the associated risk and its controls.	5 working days
Materialised risk	Management as an incident (document 37) and reassessment.	According to severity

Trigger	Action	Indicative time limit
Ineffective control on a High or Critical risk	Nonconformity and reassessment of the residual.	According to nonconformity type
Change of model, supplier, data, autonomy or regulation	Reassessment of the affected risks before the change.	Before deployment
Overdue treatment action	Alert to the owner; if High or Critical, to the committee.	On expiry

7.3 Review frequency

Residual level	In development (phases 0–5)	In production (phase 6)
Critical (exceptionally accepted)	Continuous, with a report at every committee meeting	Monthly and at every quarterly board meeting
High	At every <i>gate</i> and at least monthly	Quarterly, coinciding with Enterprise R6
Medium	At every <i>gate</i>	Half-yearly or at every R6
Low	At G3 and G5	Annually or at R6

08

Risk register

The risk register is the "Risk matrix and register" evidence of 01 §6.10. It is implemented with template **P12** and tool **T06**, on the *Risk* entity of the common data model (03 §4).

8.1 Fields

Block	Field	Content
Identification	Identifier	Initiative code and sequential number (illustrative: IA-2026-014 · R03); for portfolio risks, scope "Portfolio".
	Typical risk	Catalogue code <code>RT-<CAT>-NN</code> , or "Specific" if there is no equivalent.
	Category	EST, TEC, DAT, ECO, LEG, ORG, REP, GEN, SEG or TER.

Block	Field	Content
	Title and description	Cause → event → consequence, in plain language.
	Affected systems and suppliers	Link to the inventory (T02) and the supplier register (T09).
	Date and phase of identification	Date, phase and who identified it.
Owners	Risk owner	Named person with the ability to act.
	AI Risk Owner	Second line that issues clearance.
Inherent analysis	Likelihood (1–5) and justification	Reference used and sources.
	Impact per axis (1–5)	Economic, people and rights, regulatory, operational, reputational; determining axis.
	Inherent level	Calculated.
Controls	Controls	Description, type, code (SEG, AG or other), owner.
	Effectiveness	Effective, partially effective, ineffective or untested; date and evidence of the test.
Residual	Residual likelihood and impact	Target and verified.
	Residual level	Target and verified.
Treatment	Response	Avoid, mitigate, transfer or accept.
	Plan	Actions, owner, deadline and status (P13).
	Contingency	Trigger, actions and owner responsible for activating it.
Acceptance	Decision	Who accepts, date, validity, conditions, link to P29 if decided at a <i>gate</i> .
Monitoring	Key risk indicator	Definition, threshold, current value and date.
	Status	Identified · Under treatment · Accepted · Materialised · Closed.
	Trend and next review	Rising, stable or falling; date.
Links	Incidents, nonconformities and criteria	INC-AAAA-NNN, NC-AAAA-NNN, <i>gate</i> criteria G<n>.<nn> .

Block	Field	Content
History	Events	Assessment changes with date, author and reason (03 §2, principle 2).

8.2 Register quality

The verifier checks, as a minimum: that all Medium or higher risks have an owner, a control and a residual; that the justification of likelihood and impact is traceable; that acceptances are signed by the competent person and are current; that the typical risks applicable to the initiative's technology and classification have been considered (included or discarded with a reason); and that the register existed before the *gate* (01 §7.4, rule 3). The specific criteria are in documents 21 and 22.

09

Catalogue of typical risks

The catalogue is a starting list, not an exhaustive one. In phase 3, the team **must** review at least the typical risks in the categories that apply to its technology and exposure, and record those discarded with a reason. Each company **may** extend the catalogue with sequential codes and **must** declare those it adds.

The *Phase* column indicates where the risk is identified and where it is mainly treated; controls with a SEG or AG code are described in document 35.

9.1 Strategic (EST)

Code	Risk and description	Typical causes	Typical controls	Phase
RT-EST-01	Misalignment with the AI thesis. The initiative does not contribute to the ambition or the spheres approved in C2.	Technology-driven origin; weak sponsorship; no screening.	Fit with the thesis at G0; sphere classification (P07); prioritisation in C3.	0–1
RT-EST-02	Declared but unevidenced ambition. What is Optimise is presented as Augment or Transform, and decisions are made with wrong expectations.	Incentives to "transform"; ambiguous criteria.	Five classification questions (T05); confirmation in phase 2; review at G7.	1, 2, 7
RT-EST-03	Missed opportunity. The decision process is so slow that the advantage is lost.	Long <i>gate</i> time limits; the same return criterion for all bets.	Reference time limits (03 §3.6); criteria by ambition (01 §7.6); agility metrics.	C3, all

Code	Risk and description	Typical causes	Typical controls	Phase
RT-EST-04	Technology lock-in. The chosen platform or model is superseded and changing is costly.	Dependence on proprietary features; tightly coupled architecture.	Model abstraction layer; portability of data and prompts; review at R6.	4, 6
RT-EST-05	Withdrawal of sponsorship or funding. The initiative loses support before demonstrating value.	Changes in leadership; value not visible in time.	Stage-gated funding; learning milestones; stop criteria set in advance.	0, C3
RT-EST-06	Portfolio concentration in a single bet. A single initiative or platform concentrates investment or expected value.	Large bets without limits per stage.	Concentration limits in the appetite (13); quarterly review.	C3, C4

9.2 Technical (TEC)

Code	Risk and description	Typical causes	Typical controls	Phase
RT-TEC-01	Insufficient performance in real-world conditions. The system does not reach the threshold outside the test environment.	Validation with non-representative data; pilot without an attribution method.	Success threshold set in advance; pilot in real-world conditions (P22).	3, 5
RT-TEC-02	Degradation and drift. Performance falls because of changes in the data, the environment or user behaviour.	Dynamic environment; no monitoring.	Drift and performance monitoring (P25); retraining or retirement thresholds.	6
RT-TEC-03	Lack of robustness. Erroneous outputs in response to atypical inputs, noise or edge cases.	Testing limited to common cases.	Robustness and edge-case testing; referral to a person when out of domain.	4, 5
RT-TEC-04	Confabulation. Generative AI produces false content presented as true.	Lack of authoritative sources; out-of-scope tasks.	Retrieval over authoritative sources; citations; periodic evaluations; user notice.	4–6
RT-TEC-05	Outputs that are neither explainable nor traceable. It is not possible to reconstruct why an output was produced.	Insufficient logs; opaque models without documentation.	Logging of inputs, versions and outputs; model documentation; explanations proportionate to the use.	4
RT-TEC-06	Insufficient integration and capacity. Latency, availability or volume cannot support the process.	Testing without real load; external dependencies.	Load testing; service level agreements; graceful degradation.	4, 5

Code	Risk and description	Typical causes	Typical controls	Phase
RT-TEC-07	Unmanaged change. Changes to the model, prompts, parameters or configuration without prior evaluation.	Frequent changes; no version management.	Change management with regression evaluations; prompt versioning.	5, 6
RT-TEC-08	Irreversibility. It is not possible to return to the previous state or to the process without AI.	Rollback plan non-existent or untested; manual process dismantled.	Rollback plan tested before G5 (P19); fallback process maintained.	4, 5

9.3 Data (DAT)

Code	Risk and description	Typical causes	Typical controls	Phase
RT-DAT-01	Insufficient quality. Incomplete, erroneous or outdated data.	No control at source; legacy systems.	Quality profiling in phase 3; validation rules; data owner.	3, 4
RT-DAT-02	Non-representative data and bias. The data do not reflect the population or reproduce historical discrimination.	Biased samples; proxy variables for protected characteristics.	Representativeness analysis; fairness metrics; testing by group.	3, 5
RT-DAT-03	Lack of legal basis or incompatible purpose. Use of personal data without a legal basis or for a different purpose.	Reuse of operational data without analysis.	Legal analysis; data protection impact assessment; minimisation.	3
RT-DAT-04	Unknown lineage. It is not known where the data come from, what transformations they have undergone or under what licence.	Third-party data; manual processes.	Data and model lineage (P16); verified licences.	4
RT-DAT-05	Data exposure in the AI lifecycle. Personal or confidential data in training, context, prompts or logs.	Full unmasked logs; excessive context.	Minimisation and masking; limited log retention; SEG-07.	4, 6
RT-DAT-06	Obsolete or contradictory knowledge. The document base that feeds generative AI is outdated or has no owner.	No document governance; duplicates.	Owner and validity date per source; automatic removal of expired content.	4, 6

9.4 Economic (ECO)

Code	Risk and description	Typical causes	Typical controls	Phase
RT-ECO-01	Build cost overrun. The cost exceeds what was approved.	Open scope; underestimated integration.	Estimation with the cost categories (42); stage-gated funding; change control.	3, 5
RT-ECO-02	Rising or unpredictable recurring cost. Consumption of models, compute or licences grows faster than value.	Usage-based pricing; agent loops; unforeseen adoption.	Consumption budgets and limits; cost per transaction monitored (T13); SEG-10.	3, 6
RT-ECO-03	Unrealised value. Released capacity is not converted into savings or reallocated.	No realisation plan; organisational resistance.	Measurement rule 3 (00 §6); realisation plan at G5; monitoring (P28).	5–7
RT-ECO-04	Inflated value or double counting. The same euro is attributed to several use cases or declared without a formula.	Pressure for results; no independent validation.	Measurement rules 1, 2 and 5; validation by management control.	2, 6, 7
RT-ECO-05	Unforeseen exit cost. Replacing the supplier or retiring the system costs more than estimated.	Contracts without exit clauses; data in closed formats.	Exit plan (document 36); exit cost in the economic assessment.	3, 7

9.5 Legal and compliance (LEG)

Code	Risk and description	Typical causes	Typical controls	Phase
RT-LEG-01	Incorrect regulatory classification. A high-risk system is treated as minimal-risk or out of scope.	Classification without legal judgement; change of purpose.	Classification with T07 and legal review; review at R6 and upon changes (34).	0, 3, 6
RT-LEG-02	Prohibited practice. The use falls within a practice prohibited by the AI Act .	Lack of awareness; evolution of the use.	Screening at G0 and G3; always avoid (01 §6.5).	0, 3
RT-LEG-03	Transparency breach. People do not know they are interacting with an AI or that content is synthetic.	Design without transparency requirements.	Requirements of Article 50 of the AI Act in the design; interface review.	4, 5
RT-LEG-04	Automated decisions without safeguards. Decisions with significant effects without effective human intervention or the right to contest.	Nominal oversight; progressive automation not reviewed.	Human oversight design (P17); analysis of Article 22 of the GDPR .	3, 4

Code	Risk and description	Typical causes	Typical controls	Phase
RT-LEG-05	Omitted impact assessments. The data protection impact assessment or the fundamental rights impact assessment is not carried out or not updated.	Lack of awareness of the requirement; deadlines.	Classification (P11); assessments as G3 evidence.	3, 6
RT-LEG-06	Intellectual property infringement. Training data or outputs infringe third-party rights, or ownership of the outputs is unclear.	Unlicensed data; supplier terms not reviewed.	Licence review; ownership and indemnity clauses (36).	3, 4
RT-LEG-07	Late incident reporting. A serious incident, a data breach or a DORA or NIS2 incident is not reported in time.	No severity criteria; no reporting owner.	Response plan (P26); reporting matrix (37).	6
RT-LEG-08	Breach of sector-specific or resilience obligations. Requirements of sector supervisors, DORA or other regulations not carried over into the system.	Incomplete regulatory mapping.	Regulatory mapping (34); context statement (P02).	0, 3

9.6 Organisational (ORG)

Code	Risk and description	Typical causes	Typical controls	Phase
RT-ORG-01	Lack of adoption. Users do not use the system or use it only marginally.	Poor integration into work; distrust; insufficient training.	Adoption plan (P20); adoption target at G2 and G5.	4–6
RT-ORG-02	Dependence on key people. Knowledge of the system lies with a few people or with the supplier.	Scarce documentation; small teams.	Documentation and operations manual (P24); deputies; knowledge transfer.	4, 6
RT-ORG-03	Unmanaged effect on people. Changes in role, workload or employment without planning or dialogue.	Efficiency-only view; late communication.	People plan (50); consultation with employee representatives where appropriate.	2, 4
RT-ORG-04	Ineffective human oversight. People approve out of routine or without sufficient information.	High volume of approvals; automation bias.	Oversight design with sampling, time allowances and rotation; measurement of the rejection rate.	4, 6
RT-ORG-05	Unauthorised use of AI (<i>shadow AI</i>). Employees use unapproved tools with company data.	Lack of approved alternatives; policy not known.	Acceptable use policy (31); corporate alternatives; detection (T21).	C1, C4

Code	Risk and description	Typical causes	Typical controls	Phase
RT-ORG-06	Roles without segregation of duties. Those who build also verify, accept risks or decide.	Small teams; roles not assigned.	Role register (P03); incompatibilities (01 §8.2).	0
RT-ORG-07	Failure to inform and consult. Workers' representatives or affected workers do not receive, before use, the required information or consultation on a system that affects working conditions or employment.	Late or missing legal analysis; changes in parameters, rules or purpose not communicated.	Information and consultation process (50 §7.2) and information sheet (50 §7.3); no <i>Proceed</i> or <i>Proceed with conditions</i> at G5 without evidence that information has been provided; IND-ADO-12 (41).	3–5, 6
RT-ORG-08	Labour dispute. The introduction of AI gives rise to collective consultations, complaints or disputes that delay, condition or block the initiative.	Late or ambiguous communication; effect on employment without dialogue; destinations of released capacity decided without participation.	Position on AI and work approved in C2 (13); effect assessment (50 §3.4); dialogue with workers' representatives and internal communication (50 §7 and §10); senior management decision on relevant PER-D1 destinations.	2–5, 7
RT-ORG-09	Deterioration of wellbeing. Work intensification, excessive monitoring, loss of autonomy or oversight workload harm the health and engagement of the people affected.	Objectives and workloads not reviewed after AI; usage data used to monitor people; excessive review volumes.	Review of the occupational risk assessment, including psychosocial risks, and purpose limitation (50 §9.2); oversight design (P17); pulse survey (IND-ADO-15, 41).	4–6

9.7 Reputational (REP)

Code	Risk and description	Typical causes	Typical controls	Phase
RT-REP-01	Discriminatory outcomes. The system treats groups of people less favourably.	Bias in data or design; no testing by group.	Bias testing before G5; monitoring by group; human oversight.	3, 5, 6
RT-REP-02	Errors visible to customers. Wrong responses or decisions with public repercussions.	Direct exposure without safeguards.	Prior evaluations; referral to a person; rapid incident response.	5, 6
RT-REP-03	Use perceived as inappropriate or intrusive. Use that is legal but contrary to the expectations of customers, employees or society.	No ethical analysis; insufficient communication.	Ethical review in phase 3; transparency; complaints channel.	3
RT-REP-04	Harmful or offensive content. Generative AI produces content that is offensive, dangerous or unbecoming of the brand.	Insufficient filters; user manipulation.	Input and output filters (SEG-03); adversarial testing; sample oversight.	4–6

Code	Risk and description	Typical causes	Typical controls	Phase
RT-REP-05	Misleading communication about AI. Capabilities or results that AI does not have are attributed to it.	Commercial pressure; unvalidated value.	Review of communications; use of validated value (00 §6).	5, 7

9.8 Generative AI and agents (GEN)

Code	Risk and description	Typical causes	Typical controls	Phase
RT-GEN-01	Direct prompt injection and <i>jailbreak</i>. A user manipulates the system to bypass its restrictions.	Restrictions based solely on instructions to the model.	SEG-02, SEG-03, SEG-11; limits enforced outside the model.	4, 5
RT-GEN-02	Indirect injection. Malicious instructions in emails, documents, websites or tool responses alter the task.	The system treats untrusted content as instructions.	SEG-02; AG-12; AG-05; specific testing (AG-18).	4, 5
RT-GEN-03	Excessive agency. The system has more tools, permissions or autonomy than necessary.	Inherited permissions; unjustified autonomy level.	AG-02; autonomy level decided in phase 4 (35 §5); AG-20.	4, 6
RT-GEN-04	Unauthorised or untraceable action. An agent executes an action that does not correspond to an authorised intent, or it is not possible to reconstruct why.	No intent-based access control; incomplete logs.	AG-04, AG-05, AG-06, AG-10.	4–6
RT-GEN-05	Information leakage. Personal or confidential data or system prompts leak through responses, retrieval or tools.	Retrieval without user permissions; secrets in prompts.	SEG-05, SEG-06, SEG-07.	4, 6
RT-GEN-06	Improper output handling. The model's output is executed or inserted into other systems without validation.	Trust in the output; direct integrations.	SEG-04; AG-11.	4, 5
RT-GEN-07	Memory or context poisoning. Manipulated content persists in memory or in the knowledge base and alters future decisions.	Shared memory without control; ingestion without validation.	AG-14; SEG-08; RT-DAT-06.	4, 6
RT-GEN-08	Cascading failures and unbounded consumption. Loops, retries or chains of agents amplify errors or costs.	No iteration limits or circuit breakers.	AG-07, AG-16; SEG-10.	4, 6
RT-GEN-09	Human trust exploitation. The system induces a person to approve an inappropriate action.	Approvals without context; overreliance.	AG-08 with sufficient information; RT-ORG-04.	4, 6

9.9 Security and offensive AI (SEG)

Code	Risk and description	Typical causes	Typical controls	Phase
RT-SEG-01	Impersonation with synthetic voice or video. CEO fraud or false orders with a legitimate appearance.	Payment processes based on voice or image recognition.	SEG-16, SEG-17, SEG-18.	C4, 6
RT-SEG-02	AI-generated phishing and social engineering. Personalised, convincing messages at scale.	Weak authentication; public exposure of employee information.	SEG-15, SEG-17.	C4
RT-SEG-03	Accelerated exploitation of vulnerabilities. The time between the disclosure of a vulnerability and its exploitation shrinks.	Slow patching; unknown exposed attack surface.	SEG-13, SEG-19.	6, C4
RT-SEG-04	Compromise of non-human identities. Theft or abuse of agent credentials, service keys or <i>tokens</i> .	Long-lived credentials; excessive permissions.	AG-01, AG-02, AG-03, AG-20.	4, 6
RT-SEG-05	Data or model poisoning. Manipulation of training, fine-tuning or evaluation data.	Open sources; uncontrolled ingestion.	SEG-08; lineage (P16).	4, 5
RT-SEG-06	Model extraction and service abuse. Mass queries to replicate the model, infer data or exhaust resources.	Exposed interfaces without limits.	SEG-10, SEG-12.	4, 6
RT-SEG-07	AI supply chain compromise. Malicious or vulnerable models, libraries, connectors or tool servers.	Components downloaded without verification.	SEG-09; AG-13; supplier assessment (36).	4, 6

9.10 Third parties (TER)

Code	Risk and description	Typical causes	Typical controls	Phase
RT-TER-01	Supplier dependency and lock-in. Changing supplier is unfeasible within a reasonable time or cost.	Closed formats; proprietary features.	Exit strategy; portability (36 §7).	3, 4
RT-TER-02	Use of data by the supplier. The supplier uses data, prompts or outputs for training or other purposes.	Non-negotiated standard terms.	No-training-use clause; configuration verification (36 §6).	3
RT-TER-03	Unilateral change of model or terms. The supplier changes the version, behaviour or price, or withdraws the model.	Managed services without contractual advance notice.	Advance notice and pinned versions; regression evaluations (RT-TEC-07).	3, 6

Code	Risk and description	Typical causes	Typical controls	Phase
RT-TER-04	Uncontrolled sub-processors and location. Data processed by third parties or in unauthorised locations.	Opaque subcontracting chain.	List of sub-processors; authorisation of changes; contractual location.	3
RT-TER-05	Supplier failure or discontinuity. Prolonged unavailability, insolvency or market exit.	Small or critical supplier with no alternative.	Continuity and exit plan; financial due diligence; level N3 (36 §4).	3, 6
RT-TER-06	Unidentified embedded AI. Software already contracted incorporates AI features without assessment.	Supplier updates; activation by default.	Embedded AI questionnaire; review of updates (36 §8).	C1, 6
RT-TER-07	Concentration in common suppliers. Many systems depend on the same model or cloud provider.	Standardisation without concentration analysis.	Concentration analysis (section 10); limits in the appetite.	C3

10

Portfolio risk

The risks of initiatives do not simply add up: they interact. The AI Committee **must** review the portfolio, at least quarterly, with three analyses.

10.1 Concentration

Factor	What is measured	Indicative warning signal
Model provider	Systems in production and expected value that depend on the same provider or model family.	More than half of validated value depends on a single provider without a tested exit plan.
Platform and cloud	Systems on the same execution or agent platform.	Critical function with no execution alternative.
Critical function	Critical or important functions supported by AI.	More than one critical function without an operational fallback process.
Data	Systems that depend on the same data source or knowledge base.	Source with no owner or with unmeasured quality.
People	Systems that depend on the same key people.	One person is the AI Technical Owner of several Enterprise systems.
Sphere and ambition	Distribution of High residual risk by sphere and ambition level.	Risk concentrated in customer or decision spheres without enhanced oversight.

The thresholds are set in the risk appetite (document 13); those in this table are indicative.

10.2 Correlation

Two risks are correlated when the same cause materialises both at once. Typical common causes: the same foundation model (a version change or a vulnerability affects all the systems that use it), the same attack technique (an indirect injection that works on one agent usually works on others with the same design), the same regulatory change or the same legal interpretation, and the same supplier or sub-processor.

For each common cause, the AI Office **should** build a **portfolio scenario** and assess it with the matrix in section 4, adding up the economic impacts of the affected systems and taking the worst impact on the other axes. Minimum recommended scenarios:

1. Unavailability of the main model provider for several days.
2. Change in the behaviour of a foundation model used by several systems.
3. Exploited vulnerability in a common agent component (connector, tool server, library).
4. Regulatory reclassification of a type of use present in several initiatives.
5. Data breach in a shared knowledge base.

10.3 Supplier dependency

Dependency is assessed using the N1–N3 requirement level of document 36 and two questions: how much time and cost it takes to replace the supplier, and which functions stop in the meantime. Portfolio risks are recorded in T06 with scope "Portfolio", an owner in the AI Office and acceptance by the AI Committee or, if they are Critical, by the board.

11

Risk appetite and key risk indicators

11.1 Relationship with the appetite

The risk appetite is approved in C2 and developed in document 13. This methodology requires it to include, as a minimum:

- A statement per risk category (for example, zero tolerance for prohibited practices and for untreated discrimination).
- The economic materiality threshold M of section 4.2.
- Portfolio limits: maximum number of accepted High residual risks, no Critical risk without board approval, concentration limits.

- Key risk indicator thresholds.

11.2 Key risk indicators

The thresholds are indicative and are approved in C2. All are calculated with data from T01, T06, T08 and T09.

Indicator	Definition	Indicative threshold
Critical residuals	Current Critical residual risks, with and without board approval.	0 without approval
High residuals	Number and evolution of accepted High residual risks.	No increase for two consecutive quarters
Expired acceptances	Acceptances whose validity has expired without review.	0
Untested controls	Percentage of controls on High or Critical risks without a current effectiveness test.	Less than 10 %
Overdue actions	Treatment actions for High or Critical risks that are past their deadline.	0 overdue by more than 30 days
Materialised risks	Incidents linked to registered risks versus incidents with no previously identified risk.	Downward trend in unidentified ones
Drift	Systems with drift or performance metrics outside threshold.	0 without an open action
Injection success	Percentage of adversarial injection tests that succeeded in the last campaign.	Downward trend; 0 on sensitive actions
Agents with excessive permissions	Agent identities with unjustified permissions in the last review.	0
Supplier concentration	Proportion of validated value dependent on the main model provider.	According to the C2 limit
Overdue classification	Systems whose regulatory classification has not been reviewed within the time limit.	0
Unauthorised use	<i>Shadow AI</i> cases detected and not regularised.	Downward trend

Risk reporting

Recipient	Frequency	Minimum content
Body that decides the gate	At every <i>gate</i>	Inherent and residual profile, High and Critical risks with controls and effectiveness, required acceptances, changes since the previous <i>gate</i> .
AI Committee	Monthly	Residual heat map of the portfolio, new and escalated risks, pending and expired acceptances, indicators outside threshold, materialised risks.
Board or board committee	Quarterly	Critical and High residuals by sphere, concentration and portfolio scenarios, compliance with the appetite, significant incidents, exposure to offensive AI (35), decisions requested. Presented with the quarterly second-line report (P42).
C5 · Annual review	Annual	Effectiveness of the risk management system, recalibration of scales and catalogue, lessons learned.

The report to the board uses business language: what can happen, whom it affects, what is being done and what decision is requested. It should not present unprioritised technical lists of risks.

Associated tools and templates

Code	Name	Use in this document
P12	Risk matrix and register	Register with the fields in section 8, inherent and residual heat map. Phases 3, 5 and 6.
P13	Mitigation and contingency plan	Treatment actions, target residual, contingencies and triggers (sections 5–7). Phase 3.
P42	Quarterly second-line report	Quarterly portfolio risk reporting to the board committee (sections 10–12).
T06	Risk matrix and register	Calculation of levels with the extreme impact rule, inherent and residual 5 × 5 matrix for each initiative and for the portfolio, control effectiveness, acceptance by the body for its level, observations from sections 4 to 8 (Critical residual without approval, High without contingency, acceptance expired or insufficient, review overdue) and CSV export. "Risks" view of the T01 register. Portfolio concentration and correlation (section 10) are analysed with P12 §7.

Code	Name	Use in this document
T01 · T02 · T08 · T09 · T17	Initiative register, inventory, incidents, suppliers and board dashboard	Sources and destinations of risk data.

14

Related documents

Document	Relationship
01 · Foundational methodology	Section 10 (risks) and <i>gate</i> rules that this document develops.
13 · AI thesis, ambition and risk appetite	Appetite, tolerances, materiality threshold and indicator thresholds.
14 · Portfolio management	Use of portfolio risk in prioritisation and retirement.
21 · Gate and audit criteria · 22 · Checklists	G3, G4, G5 and R6 criteria on the risk register.
34 · Regulatory mapping	Obligations that give rise to risks in the LEG category.
35 · AI and agent security	SEG and AG controls cited in the catalogue.
36 · AI third parties and suppliers	N1–N3 levels and controls for TER risks.
37 · Nonconformities and incidents	Treatment of materialised risks and ineffective controls.
42 · AI costs	Cost categories for ECO risks.

15

Version control

Version	Date	Changes
0.1	16- 09- 2026	First version. Sets the common scales of likelihood and impact across five axes, the 5 × 5 matrix, target and verified residual risk, control effectiveness, tiered acceptance, the register fields, the catalogue of 70 typical risks in ten categories, the portfolio risk analysis and the key risk indicators. Organisational risks RT-ORG-07 (information and consultation), RT-ORG-08 (labour dispute) and RT-ORG-09 (wellbeing), proposed in document 50, added to the catalogue.

External sources cited

Links to the official or primary source, checked before publication. Always verify the version in force on the date of use.

Directive (EU) 2022/2555 of the European Parliament and of the Council of 14 December 2022 on measures for a high common level of cybersecurity across the Union (NIS 2 Directive) · <https://eur-lex.europa.eu/eli/dir/2022/2555/oj/eng> · Parlamento Europeo y Consejo de la UE · 14-12-2022

Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation) · <https://eur-lex.europa.eu/eli/reg/2016/679/oj/eng> · Parlamento Europeo y Consejo de la UE · 27-04-2016

Regulation (EU) 2022/2554 of the European Parliament and of the Council of 14 December 2022 on digital operational resilience for the financial sector (DORA) · <https://eur-lex.europa.eu/eli/reg/2022/2554/oj/eng> · Parlamento Europeo y Consejo de la UE · 14-12-2022

Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act) · <https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng> · Parlamento Europeo y Consejo de la UE · 13-06-2024

Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1) · <https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-generative-artificial-intelligence> · National Institute of Standards and Technology (NIST) · 26-07-2024

ISO 31000:2018 Risk management — Guidelines · <https://www.iso.org/standard/65694.html> · ISO (ISO/TC 262) · 14-02-2018

ISO/IEC 23894:2023 Information technology — Artificial intelligence — Guidance on risk management · <https://www.iso.org/standard/77304.html> · ISO/IEC (JTC 1/SC 42) · 06-02-2023

Artificial Intelligence Risk Management Framework (AI RMF 1.0), NIST AI 100-1 · <https://www.nist.gov/publications/artificial-intelligence-risk-management-framework-ai-rmf-10> · National Institute of Standards and Technology (NIST), Departamento de Comercio de EE. UU. · 26-01-2023
