

DOCUMENT 52 · AI OPERATIONS MANUAL

AI operations manual

Operation of predictive ML, generative AI, agents and third-party AI:
monitoring, changes, incidents, continuity, records and retirement

Fernando García · SEACHAD Version 0.1 (working draft) 16-09-2026 Draft for review

4

system types with
differentiated
operation

9

monitoring layers

3

change classes with
a gate rule

4

operations
templates P24–P27

IN THIS DOCUMENT

01 Purpose and scope

02 Operating principles

03 Operating model

04 Operation by type of system

05 Monitoring and alerts (P25)

06 Change management

07 Incident management

08 R6 continuity review

09 Post-market monitoring

10 Continuity and rollback

11 Records and traceability

12 Technical retirement

13 Operations indicators

14 Lite and Enterprise operation

15 Associated tools and templates

16 Related documents

17 Version control

Legal notice and disclaimer. SEVEN-G is a reference methodological framework provided "as is" and for information purposes only. It does not constitute legal, regulatory, financial or professional advice, nor does it guarantee compliance with any law or standard. References to general regulation (such as the EU AI Act, the GDPR, DORA or NIS2), to technical standards and to regulation specific to each sector or jurisdiction may be incomplete, may not apply to a particular case or may become out of date as a result of regulatory changes, interpretations or supervisory positions after their consultation date. **Each organisation that uses SEVEN-G is solely responsible for identifying the regulation that applies to it, verifying that it is current and certifying its own regulatory compliance**, with the appropriate qualified advice. The author and SEACHAD accept no liability for any use made of this content or for any decisions taken on the basis of it. The data, figures, companies and cases in the examples are fictitious or illustrative.

01

Purpose and scope

This document develops **phase 6 • Operation and governance** of the lifecycle (01 §6.8) and sets out how AI systems in production are operated so that they **continue to work, deliver value and remain under control**. It is generic: each system specifies what is established here in its operations manual (P24).

It covers:

- The specific operation of each type of system: **predictive ML, generative AI, agents and embedded third-party AI**.
- **Monitoring and alerts** (P25), with metrics and typical thresholds.

- **Change management** and which changes require a return to a *gate*.
- **Incident management** (P26, P27), with the S1–S4 scale, which is developed in document 37.
- Preparation of the **R6 continuity review**.
- The **post-market monitoring** required by regulation.
- **Continuity and rollback**.
- **Records and traceability**, including their retention.
- **Technical retirement**.

Audience. AI operations and technical owners, platform and security teams, risk owners and AI auditors.

It does not cover the building of solutions (document 53), security controls in detail (document 35) or the contractual management of suppliers (document 36).

This document contains regulatory references and does not constitute legal advice.

02

Operating principles

1. **An AI system can fail without going down:** quality, drift, bias, cost and actions are monitored, not just availability.
2. **No owner, no production:** every system has an operations owner who can be reached in accordance with its service level.
3. **Every alert has an owner and an action;** otherwise it is noise and is corrected.
4. **Changing means validating again,** in proportion to the effect of the change.
5. **Stopping is always possible,** with a tested procedure; agents have a kill switch.
6. **What is not recorded did not happen.**
7. **The provider changes too,** and its changes are treated as the company's own.

03

Operating model

3.1 Responsibilities

In phase 6 the **AI Operations Owner** is accountable and responsible (01 §8.4). The **technical owner** carries out retraining, evaluations and changes and is accountable for champion/challenger; the

product owner is accountable for value and adoption at R6; the **risk owner** grants clearance for significant changes at Enterprise; **security** is accountable for the review of agent permissions; the **AI Auditor** verifies tests, permissions, R6 and retirements.

3.2 Readiness to operate

G5 cannot be resolved with *Proceed* unless the following exist and have been verified: the **operations manual** (P24), the **monitoring and alerts configuration** (P25) with tested alerts, the **incident response plan** (P26) and an open **incident and change log** (P27), in addition to the **tested rollback plan** (P19). The operations owner **must** have taken part in phase 5 and formally accept the handover.

3.3 Operating requirements by autonomy level

Level	Additional minimum operating requirements
A0 · Assistance	Quality of outputs and cost; channel for reporting errors.
A1 · Recommendation	Plus logging of the recommendation and the human decision; human override rate (PER-10).
A2 · Supervised action	Plus logging of each action with its intent; technical limits; tested kill switch; sample-based review.
A3 · Autonomous action	Plus real-time anomaly detection; thresholds on aggregates; permissions reviewed quarterly; on-call staff able to shut down.

04

Operation by type of system

4.1 Predictive ML

4.1.1 DRIFT

Type	What changes	How it is detected	Difficulty
Data drift	The distribution of the input variables relative to that of training or validation.	Comparison of distributions by variable and overall (for example, population stability index or distances between distributions).	It is detected early and without labels, but not all data drift degrades the model.

Type	What changes	How it is detected	Difficulty
Prediction drift	The distribution of the model's outputs.	Comparison with the reference distribution.	A useful early signal when labels are delayed.
Concept drift	The relationship between the variables and what is predicted.	Performance against actual labels; business indicators as a proxy.	It requires labels, which may arrive weeks or months late.

Rule. For each model, P25 **must** state the **label delay** and which proxy indicators are used in the meantime. A model with no way of measuring its actual performance within a reasonable time **should** not pass G5.

4.1.2 RETRAINING

Mode	When	Conditions
Scheduled	At a fixed frequency defined in P24.	Same validated source data, variables and method; goes through champion/challenger.
Triggered	When a drift or performance alert justifies it.	The cause is analysed first: retraining does not correct an input data error or a process change.
Automatic	A pipeline that retrains and proposes without intervention.	Only if the pipeline is validated, the promotion criteria are predefined and, at Enterprise, promotion to production is approved by a person.

Every retraining: uses data that pass the quality controls (document 51), freezes and versions the dataset, updates P16 and is recorded in P27 as a change.

4.1.3 CHAMPION/CHALLENGER VALIDATION

Step	Activity	Criterion
1 • Offline evaluation	The new model (challenger) and the current one (champion) are evaluated on the same recent test set, not used in training.	Primary and secondary metrics defined in P24.
2 • Segments and bias	Comparison by relevant segments and bias testing.	The challenger does not worsen beyond the tolerance in any protected segment.
3 • Shadow mode	The challenger receives real data and produces outputs without effect.	Stability, latency, cost and agreement with the champion within expectations.
4 • Phased deployment	Part of the traffic is moved to the challenger, or a controlled test when the design allows it.	No critical alerts during the defined period.

Step	Activity	Criterion
5 • Decision	Promotion, rejection or extension of the test.	Promotion criteria set before the evaluation; decision-maker recorded.
6 • Retention	The previous champion remains available for rollback.	For the period defined in P19.

At Lite, steps 3 and 4 may be omitted if the system is A0 and has no direct exposure, with justification in P27.

4.2 Generative AI

4.2.1 CONTINUOUS EVALUATIONS

- **Reference evaluation set:** representative, difficult, sensitive and adversarial cases (including prompt injection), versioned and expanded with each incident and failure detected.
- **Full regression evaluation** for any change of model, instructions, sources, filters or parameters; without a satisfactory result, the change is not promoted.
- **Periodic evaluation** even when there are no changes by the company, to detect silent changes by the provider.
- **Production sampling** with human review; size and frequency in P25.
- **Evaluator model** only if its agreement with human evaluators has been measured and it is recalibrated periodically.

4.2.2 RESPONSE QUALITY AND HALLUCINATIONS

Dimension	Typical metric
Faithfulness to the sources or data provided	Responses without unsupported statements (OPE-09; document 51, CNC-04).
Correctness	Correct responses in the reference set.
Relevance and completeness	Rating against criteria.
Appropriate refusal	Rate of unwarranted refusals and of inappropriate responses.
Format and tone	Compliance with rules.
Content safety	Findings of harmful content, personal or confidential data per thousand interactions.

A **hallucination** is a statement presented as true that is not supported by the available data or sources or is false. Its management combines design (source retrieval, mandatory citation, instructions), filters, human oversight in sensitive uses and continuous measurement. A hallucination with an effect on customers or decisions is an **incident**.

4.2.3 FILTERS

Filter	What it controls	What is monitored
Input	Prompt injection attempts, personal or confidential data, prohibited content.	Trigger rate and sudden variations in it (possible attack or change of use).
Output	Personal data, harmful content, prohibited statements (for example, contractual promises), format.	Trigger rate, reported false positives.
Retrieved context	Instructions embedded in documents or pages (indirect injection).	Detections and originating sources.

Filters are **critical controls**: deactivating them is a substantial change (section 6) and their failure is an incident.

4.2.4 CONSUMPTION-BASED COSTS

- **Unit**: cost per useful transaction (query resolved, document processed), not just consumption volume.
- Monthly **budget** per system, with alerts at 80% and 100% as an initial reference.
- **Limits** per user, session and period; maximum context size; maximum calls per task.
- **Anomalies**: rising cost per transaction with no change in use (loops, excessive context, abuse).
- **On reaching the limit**: cost-driven degradation cascade (section 10.3), defined in P24 and validated at G5; complete cut-off only if it does not affect a critical function.
- **Allocation** to the *model consumption* category (document 42) and tracking in T13.

4.2.5 PROVIDER MODEL VERSIONS

1. Pin the **exact version** of the model in production (not an alias to the "latest version") and record it in P16.
2. Record announced **deprecation dates** and start the migration with enough time for a full evaluation and champion/challenger.
3. Detect **behavioural changes without a version change** through periodic evaluation and daily control queries.
4. Review notices on **data use and processing location** and reassess in P14 (document 36).

Migration to a new version from the same provider is a **significant change**; migration to another model family or to another provider is a **substantial change** (section 6).

4.2.6 USAGE DRIFT

In generative AI and agents there are no fixed input variables or labels with which to measure drift as in predictive ML, but the system also degrades when **what it is asked changes**: new topics, new types of user, queries nobody anticipated. The system was validated at G5 with a specific scope and evaluation set; if actual use moves away from them, the evaluations no longer represent what happens in production.

Signal	What is compared	How it is measured
Distribution of topics or intents	Breakdown of queries by topic or intent against that of the use validated at G5.	Classification of queries (by rules, by a classifier or by human sampling) and distance between distributions; the same population stability index as in section 4.1.1 is applied to the categories.
Queries outside the validated scope	Proportion of queries that do not correspond to any validated topic or task.	The classification above or sampling; rate over the total for the period (OPE-19).
Shape of the inputs	Length, language, channel, document type, tools invoked by agents.	Distributions against the reference.
Retrieval without a relevant source	Queries for which the search finds no suitable sources.	Rate for the period (document 51, CNC-08).

Rule. P25 **must** set the **usage reference** (topics, intents and scope validated at G5), the way queries are classified and the thresholds. When usage drift exceeds the critical level: the evaluation set is expanded with the new cases and the system is evaluated on them; if the new use falls outside the approved purpose, the scope is restricted or it is treated as a **new use** (table 6.2, return to G3).

Why it matters. An assistant validated for product queries that starts receiving complaints or legal questions still "works" and its evaluations stay green, because they measure what was anticipated, not what is happening. Usage drift is the early signal that the actual risk is no longer the approved one and that the measured value may lie elsewhere.

4.2.7 BIAS IN RESPONSES

Bias is not exclusive to models that score or decide: a generative system can **treat people or groups differently** in the content, tone, recommendations or quality of the response.

Test	What it detects	How it is done
Counterfactual pairs	Differences in response attributable to a protected attribute (sex, age, origin, disability or others set by the impact assessment) or to a proxy for it (name, language, writing style).	Part of the reference evaluation set (4.2.1) consists of pairs of cases identical except for that attribute; the outcome of each pair is compared against the P25 criteria (decision or recommendation, amount, tone, refusal, completeness).
Quality by segment	Lower accuracy or more refusals for a group.	Metrics from 4.2.2 broken down by the segments defined in the impact assessment, in the periodic evaluation and in production sampling.
Stereotyped content	Stereotyped statements or assumptions about groups.	A specific criterion in the human review of the sample and in the output filters.

Rules.

1. The **unequal response rate** (pairs with a material difference ÷ pairs evaluated, OPE-20) is measured in every regression evaluation and in the periodic evaluation; its threshold is set in P17 alongside the human oversight thresholds.
2. It is **mandatory** when the system recommends, prioritises, drafts communications or decides about people, or has direct exposure to customers; for other systems, when the risk assessment indicates it (document 33).
3. A change of model, instructions or sources is not promoted if it worsens this rate beyond the tolerance, just like the challenger in predictive ML (4.1.3, step 2).

Why it matters. In generative AI, bias does not appear in a results column that can be counted; it appears in how things are drafted, recommended or refused. Without a test designed for it, it goes unseen, and it is precisely what a regulator, an affected customer or the media will see.

4.3 Agents

The design controls for agents (identity, permissions, intent-based access control, kill switch) are defined in document 35 and assessed with T10. This section sets out how they are operated.

4.3.1 ACTION MONITORING

Each action by an agent **must** record: the agent's own identity (not shared with people) and version; the user or process on whose behalf it acts; the **declared intent** (task objective and justification for the action); the tool, operation and data used; the result and effect; prior human validation and who gave it, if applicable; a timestamp and a correlation identifier to reconstruct the sequence.

4.3.2 LIMITS

- **Per action:** amount, number of recipients, types of data.
- **Per period:** number of actions, cumulative spend, operations on the same customer or record.
- **Scope:** permitted systems, environments and functions.
- **Execution:** maximum steps, time and retries per task.
- **Escalation:** actions that always require human validation, whatever the autonomy.

Limits are enforced **technically outside the agent** (in the tools, gateways or permissions), not only through instructions to the model. An action blocked by a limit is logged and, as per P25, generates an alert.

4.3.3 INTENT ANOMALIES

An **intent anomaly** is an action or sequence that is not consistent with the declared objective of the task.

Signal	Typical response
Tools that do not correspond to the task (possible injection or planning error).	Blocking of the action and review.
Access to data or systems out of scope.	Blocking, critical alert, security review.

Signal	Typical response
Instructions originating from processed external content (indirect injection).	Blocking and analysis of the source.
Loops or repetition of the same action.	Cut-off by execution limit.
Sudden growth in the volume of actions or cost.	Pausing of the agent and review.
Attempt to extend its permissions or modify its limits.	Shutdown of the agent; incident of at least S2.

4.3.4 KILL SWITCH

- **Levels:** per action, session, agent and set of agents of a given type.
- **Who activates it:** the operations owner, the designated human supervisor and security, **without prior approval**.
- **Effect:** stops new actions, cancels pending ones, revokes temporary credentials and leaves the systems in the consistent state defined in P24, within a target time set in P25.
- **Testing:** quarterly at Enterprise, half-yearly at Lite and after each significant or substantial change (OPE-12).
- **Reactivation:** after root cause analysis, with the approval of the operations owner and, if there was an S1 or S2 incident, of the risk owner.

4.3.5 PERMISSION REVIEW

1. Inventory of each agent's **non-human identities**, with an owner.
2. **Least privilege:** each permission is justified by a task of the agent.
3. **Short-lived credentials** wherever the platform allows it.
4. **Quarterly** review for A2–A3 and **half-yearly** for A0–A1 (OPE-13).
5. Removal of permissions **not used** within the defined period (initial reference: 90 days).
6. Any extension of permissions that enables **new effects** on third parties, money, personal data or production systems is a **substantial change**.

4.4 Embedded third-party AI

Software from a supplier with AI functions that take part in the company's processes (01 §1.2).

- **New AI functions activated by the supplier**, sometimes by default: review of release notes, verification of the configuration after each version and deactivation pending assessment where the contract allows it.

- **Model or behavioural changes without notice:** process outcome indicators, sampling and regression testing before relevant updates.
- **Lack of logs:** contractual requirement for access to logs and support in incidents (document 36).
- **Changes in data processing** (location, sub-processors, training): reassessment in P14.
- **Supplier incidents:** notification channel and time limits in P26.

The intensity of these controls depends on the requirement level for the third party: **N1 Standard** · **N2 Enhanced** · **N3 Critical** (document 36).

05

Monitoring and alerts (P25)

5.1 Monitoring layers

Layer	What is monitored	Applies to
1 • Availability and infrastructure	Availability, latency, errors, capacity.	All
2 • Input data quality	Nulls, ranges, schema, freshness, volume.	All
3 • Drift	Data, predictions and concept (4.1.1); usage: topics, intents and queries outside the validated scope (4.2.6).	Data, predictions and concept: predictive ML. Usage: generative AI and agents
4 • Model performance	Metrics against the G5 validation, by segment.	Predictive ML
5 • Response quality	Faithfulness, correctness, appropriate refusal, currency of sources; bias in responses (4.2.7).	Generative AI
6 • Actions and limits	Actions, blocks, intent anomalies, permissions.	Agents
7 • Security	Filter triggers, injection attempts, improper access.	All, with emphasis on generative AI and agents
8 • Cost	Consumption, cost per transaction, budget.	All, with emphasis on generative AI and agents
9 • Human oversight, use and value	Human override rate, adoption, value hypothesis indicators, bias.	All

5.2 Typical metrics and thresholds

The values are **illustrative and to be calibrated** for each system in phase 5 on the basis of its baseline. The severity shown is the initial severity of the incident if an effect is confirmed; the final classification follows section 7.

Layer	Metric	Warning	Critical	Initial severity if there is an effect
1	Monthly availability against the service level	Below target in the partial period	Breach of the service level	S3; S2 if it is a critical function
1	95th percentile latency	> target for 15 min	> 2 × target for 15 min	S4–S3
2	Proportion of nulls in critical variables	> 2 × 30-day average	> phase 3 quality threshold	S3
2	Age of input data	> planned update interval	> 2 × interval	S3
3	Population stability index by variable	> 0.1	> 0.25	S4–S3 (common convention in practice; to be calibrated)
3	Stability index of the distribution of topics or intents	> 0.1	> 0.25	S4–S3 (to be calibrated)
3	Queries outside the validated scope	> 2 × G5 reference	> 3 × G5 reference or > 10% of queries	S3; S2 if the new use affects people or regulated topics
4	Primary metric against G5 validation	Relative drop > 5%	Relative drop > 10%	S3; S2 if it affects decisions about people or customers
4	Adverse impact ratio between groups	< P17 threshold + margin	< P17 threshold	S2
5	Faithfulness to sources in sampling	< target – 2 points	< target – 5 points	S3; S2 with direct exposure
5	Responses with outdated content	> 0 on regulated or contractual topics	Repetition after correction	S3–S2
5	Unequal response rate in counterfactual pairs	> P17 threshold – margin	> P17 threshold	S3; S2 if there are decisions or communications about people
6	Actions blocked by limits	> 2 × 7-day average	Any attempt to exceed an amount or scope limit	S3; S2 if there is an attempt to extend permissions

Layer	Metric	Warning	Critical	Initial severity if there is an effect
6	Intent anomalies	Any unexplained anomaly	Anomaly with an effect on third parties, money, personal data or production	S2; S1 if there is material harm
7	Input filter triggers	$> 3 \times 7\text{-day average}$	Sustained attack pattern	S3–S2
7	Exposure of personal or confidential data	—	Any confirmed case	S2; S1 if it is massive or involves special categories
8	Monthly cost against budget	$\geq 80\%$	$\geq 100\%$	S4–S3
8	Cost per transaction	$> 1.5 \times 30\text{-day average}$	$> 3 \times 30\text{-day average}$	S3
8	Operation in cost-driven degraded mode	Any activation	> 5 days in the month or quality in the mode below the P25 minimum	S4; S3 if quality falls below the minimum
9	Human override rate	Close to 0% for 30 days	—	Review at R6
9	Primary value indicator	Below the hypothesis for two consecutive periods	Below the stop criterion	Review at R6; brings G7 forward

5.3 Alert design rules

1. Each alert has an **owner**, a **procedure** in P24 and an **initial severity**.
2. Critical alerts reach a person who can be contacted during the system's hours of operation.
3. They are **tested** before G5 and after each change that affects them.
4. **Noise** is measured (OPE-16) and thresholds are reviewed at each R6.
5. A deactivated or silenced alert is recorded in P27 with the reason, owner and reactivation date.

Change management

6.1 Change classes

Class	Definition	Examples	What it requires
Minor	Does not alter the behaviour of the system or its controls.	Infrastructure scaling; security patch with no functional change; correction of a dashboard.	Standard testing; approval by the operations owner; recording in P27.
Significant	Alters behaviour within the approved purpose, scope, autonomy and risks.	Retraining with the same source data and method; new version of the model from the same provider and family; change of system instructions; new knowledge sources of the same type and permissions; adjustment of thresholds within approved ranges.	Regression evaluation or champion/challenger; update of P16 and P24; risk owner clearance at Enterprise; report at the next R6; verification by the auditor by sampling.
Substantial	Alters the approved purpose, scope, autonomy, data, critical controls or risk profile.	See table 6.2.	Return to <i>gate</i> as per table 6.2.

In case of doubt between two classes, the **higher** one applies.


6.2 Substantial changes and the *gate* to return to

Change	Return to	Reason
New purpose or new use of the system	G3	Value, risk and possibly the regulatory classification change.
Extension to new groups, countries, channels or customers	G3 (or a new phase 0 if decided as scaling at G7)	Exposure, regulation and effect on people change.
Increase in the autonomy level (for example, from A1 to A2)	G4 and G5	Human oversight, limits and security change.
New agent tools or permissions with effects on third parties, money, personal data or production	G4 and G5	The surface for action changes.

Change	Return to	Reason
New categories of personal data, special categories or new sources with personal data	G3	Legal basis and impact assessment.
Change of variables in systems that make decisions about people	G4 and G5	Bias testing and information to worker representatives (document 50).
Change of provider or model family	G5 (and G3 if the requirement level for the third party changes)	Full evaluation and supplier assessment.
Deactivation or replacement of a critical control (filters, human oversight, limits)	G4 and G5	Critical controls do not admit conditions (01 §7.3).
Change that may alter the regulatory classification or the intensity	G3	New classification (T07) and intensity determination (T04).
Modification that may be considered substantial under the EU AI Act in a high-risk system	G3, with legal analysis	May require a new conformity assessment and, in certain cases, turn the deployer into a provider (Arts. 3(23) and 25; document 34).

Classification of a change

The class determines the validation; substantial changes return to a gate



Source: SEVEN-G · SEACHAD

6.3 Emergency changes

- They are permitted in order to **contain an incident**. They may be applied before their ordinary approval.
- **They may only reduce** functionality, autonomy, scope or permissions (shutdown, rollback, deactivation of a function, stricter limits). Never extend them.

- They are recorded in P27 within a maximum of 24 hours and reviewed within 5 working days, with a decision on whether they are consolidated, reverted or processed as an ordinary change.

6.4 Rules

1. A substantial change applied without returning to the corresponding *gate* is a **major nonconformity**; if it affects a critical control or a high-risk system, it is **critical** (01 §12).
2. Every change to a system that affects working conditions triggers a review of whether the information provided to worker representatives must be updated (document 50, section 7).
3. Provider changes are classified using the same rules as the company's own changes.
4. **Activating a planned degraded mode** defined in P24 and validated at G5, including switching to a fallback model under section 10.3, **is not a change**: it is an approved operating mode and is recorded in P27 as an operational event. Adding a new fallback model or modifying its conditions of use is a change: **significant** if it is from the same family and provider already assessed; **substantial** otherwise (table 6.2).

07

Incident management

The complete process, the definitive time limits and the relationship with nonconformities are set out in document 37. This section establishes what is needed to operate.

7.1 What an AI incident is

Any event that produces or may produce an undesired effect owing to the operation or use of an AI system, **even if the system is available**: erroneous outputs with an effect, bias, hallucinations with

consequences, unauthorised actions by agents, exposure of information, anomalous costs, control failures, attacks and outages.

7.2 Severity


Severity	Criterion	Illustrative examples	Indicative response times (37 §4.3)	Reports to
S1 · Critical	Possible serious incident under the EU AI Act ; serious incident under DORA or NIS2 where applicable; significant material harm to people, customers, money or rights; or loss of control of a system that acts.	Agent executing unauthorised payments or communications to customers; systematic discrimination in decisions about people; massive exposure of personal data.	Triage within 1 hour, at any time; immediate coordinator; AI Committee within 4 hours; board committee within 24 hours.	AI Committee and board committee; authorities where appropriate.
S2 · High	Impact 4 on any axis; personal data breach reportable to the authority without a high risk to individuals; repeated errors visible to customers; compromise of an agent identity contained before it has any effect; failure of a critical control without harm; relevant incident at an N3 supplier.	Erroneous responses to customers about contractual terms; kill switch that does not work in testing; attempt by an agent to extend its permissions.	Triage within 4 hours; immediate coordinator; AI Committee within 24 hours; board committee in the quarterly report.	AI Committee (24 hours); board committee (quarterly report).
S3 · Medium	Impact 3; degradation beyond threshold in a relevant process; successful prompt injection without any sensitive action; limited internal leakage with no risk to individuals.	Critical drift with human review correcting the outputs; filter failing in isolated cases.	Triage within 1 working day; planned resolution; AI Committee in the monthly report; board committee in the quarterly report.	AI Committee (monthly report); board committee (quarterly report).
S4 · Low	No effect on users or decisions.	Warning alerts; minor formatting defects.	Triage within 5 working days; analysis at R6.	Recording in P27; AI Committee (monthly report).

Severity is assigned on **detection** using the most severe criterion met (one is enough; impact axes from document 33 §4.2, criteria from 37 §4.2) and reviewed once the scope is known. In case of doubt, the higher one is assigned.

7.3 Flow

Flow of an AI incident

Contain first, report on time and always learn



Source: SEVEN-G · SEACHAD

7.4 Regulatory notifications

Consulted in September 2026; time limits and cases are detailed in documents 34 and 37.

Regulation	Case	Reference
EU AI Act	Serious incidents involving high-risk systems: notification by the provider to the market surveillance authority, with a general maximum time limit of 15 days from becoming aware and shorter time limits in the most serious cases. A deployer that detects a serious incident immediately informs the provider.	Arts. 73 and 26(5)
GDPR	Personal data breach: notification to the supervisory authority without undue delay and, where feasible, within 72 hours; communication to data subjects if there is a high risk.	Arts. 33 and 34
NIS2 (as transposed)	Significant incidents: early warning within 24 hours and notification within 72 hours.	Art. 23 of Directive (EU) 2022/2555
DORA	Major ICT-related incidents in financial entities, with the time limits in its technical standards.	Regulation (EU) 2022/2554
Sector-specific regulation	As per the context statement.	P02

For each system, P26 **must** identify which notifications may apply, who decides whether notification is required and who notifies.

7.5 Relationship with nonconformities

An incident may reveal a breach of the framework (for example, a substantial change without a *gate* or a deactivated critical control). In that case a **nonconformity** is also opened, with its own scale: minor, major or critical (01 §12). Both are recorded in T08 and linked.

08

R6 continuity review

8.1 Frequency

At least **quarterly at Enterprise** and **half-yearly at Lite** (01 §6.8). It is brought forward after an S1 or S2 incident, a substantial change or a relevant change by the provider.

8.2 R6 package

The product owner and the operations owner prepare, with the data for the period:

Block	Content
Value (P28, T12)	Realised value against the hypothesis with its status; stop criteria.
Released capacity (T20)	Hours released, realised, reassigned and pending a decision.
Cost (T13)	Actual cost against budget; cost per transaction; activations and days in cost-driven degraded mode, with the quality during the mode (OPE-18).
Stability, performance and quality (P25)	Availability, critical alerts and noise; model or response metrics; drift, including usage drift (OPE-19); bias, including in responses (OPE-20).
Incidents and changes (P27, T08)	Incidents by severity and actions; open nonconformities; changes by class and by the provider.
Agents (P25, T10)	Blocked actions, anomalies, kill switch test, permission review.
Human oversight and adoption (document 50)	Human override rate, adoption, supervisor training.
Risk and compliance (P11, P12, T04, T07)	Risks; currency of classification, intensity and impact assessments; information to worker representatives.
Continuity (P19, P14)	Latest rollback test; supplier situation.

Block	Content
Proposal	Proceed with operation, Proceed with conditions or Bring G7 forward (01 §7.3), with reasons.

The R6.<nn> criteria are coded in document 21.

8.3 Outcomes

Outcome	When
Proceed with operation	Value, stability, risk and compliance within what was approved.
Proceed with conditions	Non-critical deviations with an action, deadline and owner. Not permitted for critical controls, legal compliance or human oversight.
Bring G7 forward	Relevant deviation in value (stop criterion reached), increase in risk or a better alternative.

If at G7 it is decided to keep the system unchanged, this is recorded as *Iterate* with a return to phase 6 (21 §5.3). If an unacceptable risk is detected during the review, G7 is not awaited: the incident procedure is activated and, where appropriate, the shutdown.

Omitting an R6 is a **major nonconformity** (01 §12).

09

Post-market monitoring

The [EU AI Act](#) requires **providers** of high-risk systems to establish and document a **post-market monitoring system**, based on a plan, that actively and systematically collects and analyses data on the performance of the system throughout its lifetime in order to evaluate continuous compliance (Art. 72). **Deployers** must monitor operation in accordance with the instructions for use, inform the provider where appropriate and, if they consider that the use may present a risk, inform and suspend use (Art. 26(5)). High-risk obligations apply from 2 December 2027 (Annex III) and 2 August 2028 (Annex I), following Regulation (EU) 2026/1744; the timetable and the other amendments are addressed in document 34 (§3.1 and §3.2).

- **As a provider** (including where it becomes one through substantial modification or change of purpose): monitoring plan integrated into P24 and P25, with the data collected, their periodic analysis and the connection with risks, changes, incidents and R6.

- **As a deployer:** monitoring in accordance with the provider's instructions, a channel for informing the provider, a suspension criterion and retention of logs (P25, P26).
- **Systems that are not high-risk:** the monitoring in section 5 fulfils an equivalent function.

This document does not constitute legal advice.

10

Continuity and rollback

10.1 Degraded operating modes

Mode	When
Validated previous version (model, instructions, configuration)	Degradation after a change.
Alternative without AI (rules, manual process, previous system)	Failure of the system or the provider; S1–S2 incident.
Reduced autonomy (from A2–A3 to A1)	Intent anomalies; doubts about limits.
Reduced scope to lower-risk cases or channels	Problems in a segment.
Previously assessed alternative provider ; if not assessed, it is a substantial change	Prolonged unavailability or withdrawal of the provider.
Lower-cost fallback model , validated at G5 (section 10.3)	Consumption budget exhausted or month-end forecast above 100%; volume or price peaks.

10.2 Rules

1. The **rollback plan** (P19) is tested before G5 and **periodically** in production: at least annually at Lite and half-yearly at Enterprise, and after each substantial change (OPE-15).
2. The plan identifies the **people and capabilities** needed to operate without the system and is coordinated with document 50 (section 5.4).
3. For systems that support a **critical or important function**, the maximum recovery time and the maximum tolerable data loss are set in P24 and aligned with the business continuity plan and, where applicable, with [DORA](#).
4. **Dependence on a provider** is assessed in P14 with an exit strategy (document 36).

10.3 Cost-driven degradation cascade

Systems with variable consumption (generative AI and agents) **must** have defined in P24 what happens when consumption reaches the budget, so that the response is neither improvised nor an abrupt cut-off. Degradation follows a **cascade** of levels, from least to most restrictive:

Level	What is done	Condition for using it
N0 · Normal operation	Primary model and configuration approved at G5.	—
N1 · Optimisation without a model change	Stricter limits per user and session, shorter context, reuse of responses, batch processing of non-urgent work.	Parameters and ranges approved at G5 (threshold adjustment within range).
N2 · Fallback model	Part or all of the traffic moves to a lower-cost model.	Model evaluated before G5 with the same reference evaluation set (4.2.1), including the bias tests (4.2.7), with a result within the P25 minimum quality; version pinned and recorded in P16.
N3 · Reduced scope	The system handles only the lowest-risk or highest-value cases or channels; the rest go to people.	Allocation criterion defined in P24 and people available (document 50).
N4 · Alternative without AI	Manual process, rules or previous system.	Tested rollback plan (P19).

Rules.

1. The levels, their triggers (for example, N1 at 80% of the budget, N2 at 100%) and who activates them are set in P24; their activation is tested before G5 together with the rollback plan.
2. **No model downgrade** (N2) in uses that decide or recommend about people, in critical or important functions or in high-risk systems, unless the fallback model has been validated for that use to the same standard as the primary one; in those cases the system moves directly to N3 or N4, or a budget supplement is approved (document 42, section 8.1).
3. In N2, **quality is monitored more intensively**: production sampling at least doubled and an alert if it falls below the P25 minimum; if it does, the system moves to the next level.
4. Each activation and each return to N0 is recorded in P27 with the date, level, reason, who decides and the quality measured during the mode. The return to N0 is authorised by the operations owner when there is budget or the cause has been corrected.
5. If the system spends more than a month in N2 or a higher level, the budget or the design is wrongly sized: it is taken to the next R6 with a proposal (supplement, permanent optimisation under document 42, section 8.5, or a change of the primary model as an ordinary change).
6. No degradation may **deactivate critical controls** (filters, human oversight, agent limits, records): the cascade reduces cost, not controls.

Illustrative example (fictitious figures). Customer service assistant with a monthly budget of €12,000. On day 20 consumption reaches €9,600 (80%) and the month-end forecast is €14,400: N1 is activated (shorter context and reuse of frequent responses) and the forecast falls to €13,100. On day 26 consumption reaches 100%: informational queries move to N2, with a fallback model validated at G5 that costs 70% less, while complaints remain on the primary model. Sampling on those days gives an accuracy of 91% against the P25 minimum of 88%. The month closes at €12,600 (105%), with six days in N1 and five in N2, and the R6 decides to adjust the budget to the actual growth in volume.

Why it matters. Without a defined cascade, when the budget runs out only two bad options remain: cut the service or keep spending without control. With it, cost is contained without improvisation and without touching the controls, the board knows what quality is delivered at each level, and the organisation learns whether the budget or the design was wrongly sized.

11

Records and traceability

11.1 What is recorded

- **All systems:** inputs (or their reference) and outputs with the model and instruction version and a timestamp, with minimisation; alerts and action taken; changes and incidents (P27, T08); results of evaluations and of champion/challenger.
- **A1–A3 and decisions about people:** recommendation, human decision, who decides and the reason for any override.
- **A2–A3 agents:** the fields in section 4.3.1.
- **Generative AI:** sources and versions retrieved for each response; filter triggers and blocks.

11.2 Retention

Reference	What it establishes	Rule in SEVEN-G
EU AI Act , Arts. 12, 19 and 26(6)	High-risk systems allow for the automatic recording of events. Providers and deployers keep the automatically generated logs under their control for a period appropriate to the intended purpose, of at least six months , unless otherwise provided by other applicable law.	Minimum of six months for high-risk systems, extendable according to the purpose, sector-specific regulation and claim periods.
EU AI Act, Art. 18	Providers of high-risk systems keep the technical documentation and other documentation for ten years after the system has been placed on the market or put into service.	Applies when the company is a provider.
GDPR , Art. 5(1) (c) and (e)	Data minimisation and storage limitation.	Logs containing personal data are limited to what is necessary, with pseudonymisation where possible and a defined retention period.
Sector-specific and commercial regulation	Their own time limits (for example, in financial services).	Incorporated into the context statement (P02).
Claim and limitation periods	Periods during which it may be necessary to substantiate a decision.	Considered with legal counsel for decisions about people and customers.

Integrity rules. Records are protected against modification, with restricted and traced access; their deletion when the period expires is automatic and logged. In the event of an incident, complaint or investigation, **deletion is blocked** for the records concerned.

This document does not constitute legal advice.

12

Technical retirement

The decision to retire is taken at G7 and documented in P30; its management is supported by T22. This section sets out the technical execution.

- 1. Preparation:** replacement or alternative without AI operational, people trained and communication carried out (P30).
- 2. Freeze:** no changes except emergencies; capture of the final state and metrics (P27).
- 3. Progressive disconnection:** removal of traffic and of the integrations and processes that consume outputs.

4. **Agents:** kill switch, cancellation of tasks and **revocation of all credentials and non-human identities**, with logging.
5. **Data:** deletion, anonymisation or archiving as per G7 and the retention periods (document 51, section 4.5), with a record of deletion.
6. **Models, instructions and indexes:** archiving of the final version where appropriate; deletion of indexes derived from sources containing personal data; P16 with status *retired*.
7. **Records:** retained until the end of their period (section 11.2); never deleted with the retirement.
8. **Suppliers:** termination of services, confirmation of data deletion by the supplier, closure of licences and consumption (document 36).
9. **Inventory and register:** status *Retired* in T01 and T02 with date, coded reason, deciding body and replacement.
10. **Subsequent verification:** after a defined period, check that no access, costs or dependent processes remain; closure note verified by the AI Auditor at Enterprise.

13

Operations indicators

The codes are provisional until they are consolidated in document 41. The thresholds are **indicative and to be calibrated**.

Code	Indicator	Formula	Frequency	Indicative reference
OPE-01	Availability	Time available ÷ time committed in the service level	Monthly	As per service level
OPE-02	Monitoring coverage	Systems in production with a current P25 and tested alerts ÷ systems in production	Quarterly	100%
OPE-03	Mean time to detect	Σ (time of detection – time of onset) ÷ number of incidents	Quarterly	Downward trend
OPE-04	Mean time to contain	Σ (time of containment – time of detection) ÷ number of incidents, by severity	Quarterly	Within the references in 7.2
OPE-05	Incidents by severity	S1, S2, S3 and S4 incidents ÷ systems in production, per quarter	Quarterly	Trend
OPE-06	Substantial changes without a <i>gate</i>	Substantial changes applied without returning to the <i>gate</i> ÷ substantial changes	Quarterly	0%

Code	Indicator	Formula	Frequency	Indicative reference
OPE-07	Failed change rate	Changes reverted or giving rise to an incident ÷ changes applied	Quarterly	Downward trend
OPE-08	Untreated drift	Models with drift at critical level with no action recorded within the P25 time limit ÷ monitored models	Monthly	0%
OPE-09	Responses with unsupported statements	Evaluated responses with at least one unsupported statement ÷ evaluated responses	Monthly	Threshold per system
OPE-10	Cost per transaction and variance	Operating cost for the period ÷ useful transactions; and actual cost ÷ budget	Monthly	As per hypothesis and budget
OPE-11	Blocked agent actions	Actions blocked by limits ÷ actions attempted; and intent anomalies reviewed ÷ detected	Monthly	100% of anomalies reviewed
OPE-12	Tested kill switch	A2–A3 agents with a kill switch test on time and passed ÷ A2–A3 agents	Quarterly	100%
OPE-13	Reviewed agent permissions	Agent identities reviewed on time ÷ agent identities	Quarterly	100%
OPE-14	R6 on time	Continuity reviews carried out on time ÷ reviews due	Quarterly	100%
OPE-15	Tested rollback	Systems with a rollback test on time ÷ systems in production	Half-yearly	100%
OPE-16	Alert noise	Alerts closed without action ÷ alerts generated	Monthly	Downward trend
OPE-17	Supplier versions at risk	Systems whose provider model version has an announced deprecation with no migration plan ÷ systems with a provider model	Monthly	0%
OPE-18	Cost-driven degraded mode	Days in the period at N1 or above in the cascade (10.3), by level; and quality measured in the mode ÷ P25 minimum quality	Monthly	Downward trend; quality ≥ 100% of the minimum
OPE-19	Queries outside the validated scope	Queries classified outside the topics or tasks validated at G5 ÷ queries in the period	Monthly	According to the G5 reference

Code	Indicator	Formula	Frequency	Indicative reference
OPE-20	Unequal responses in counterfactual pairs	Pairs with a material difference in outcome ÷ pairs evaluated	At each regression and periodic evaluation	P17 threshold

14

Lite and Enterprise operation

Aspect	Lite	Enterprise
Operations manual (P24)	Simplified	Complete
Monitoring layers	1, 2, 5 or 4 depending on type, 8 and 9	All applicable
Continuous evaluation of generative AI	Periodic and upon changes	Periodic, upon changes and with production sampling
Champion/challenger	Steps 1, 2, 5 and 6	Complete
Significant changes	Approval by the operations owner	Risk owner clearance
On-call	Working hours, except A2–A3	According to the system's hours of operation
Kill switch test	Half-yearly	Quarterly
Agent permission review	Half-yearly	Quarterly
Rollback test	Annual	Half-yearly
Cost-driven degradation (10.3)	At least N1 and N4 defined; N2 optional	Full cascade defined and tested before G5
Usage drift (4.2.6)	Queries outside the scope by sampling	Distribution of topics and queries outside the scope with systematic classification
Bias in responses (4.2.7)	When there are decisions or communications about people or direct exposure	The same, and whenever the risk assessment indicates it
R6	Half-yearly	Quarterly
Verification	AI Office; auditor by sampling	AI Auditor

Associated tools and templates

15.1 Operations templates

Code	Template	Minimum content	Reviewed at
P24	Operations manual	Description of the system and type; autonomy level; service levels; responsibilities and on-call; procedures per alert; retraining or continuous evaluation; agent limits and shutdown; degraded modes and cost-driven degradation cascade with its triggers; costs and consumption limits; records and retention; dependencies and suppliers; post-market monitoring plan where applicable; maximum recovery time for critical functions.	G5, R6, after significant and substantial changes
P25	Monitoring and alerts configuration	Metrics per layer; baseline; warning and critical thresholds; initial severity; owner; procedure; label delay; usage reference and query classification; reference evaluation set with counterfactual pairs and sampling; minimum quality in degraded mode; result of the alert test.	G5, R6
P26	Incident response plan	S1–S4 criteria applied to the system; contacts and escalation; pre-approved containment actions; applicable notifications with owner and time limit; communication to users, affected parties and the supplier; relationship with the continuity plan.	G5, R6, after each S1–S2 incident
P27	Incident and change log	Changes: class, description, validation, approval, <i>gate</i> where applicable, date. Incidents: severity, timeline, containment, notifications, root cause, actions, closure, link to nonconformity.	Continuous; R6

15.2 Other tools and templates

- **Tools:** T08 (incidents and nonconformities, supporting P27), T10 (agent security), T13 (operating costs), T12 (value at R6), T22 (retirements), T02, T04 and T07 (currency of inventory, intensity and classification).
- **Templates:** P16 (lineage, updated with each change), P19 (rollback), P28 (value at R6), P29 (R6 outcome and returns to *gate*), P30 (retirement).

Related documents

Document	Relationship
01 and 03 · Methodology and register	Phase 6, R6, <i>gate</i> rules, nonconformities; statuses and events.
21 and 33 · Gate criteria and risks	R6 and G5 criteria; risk register in operation.
34 · Regulatory mapping	Post-market monitoring, logs, serious incidents and substantial modification.
35 and 36 · Security and third parties	Design of agent controls; N1–N3 levels and exit strategy.
37 · Nonconformities and incidents	Complete process, S1–S4 scale and definitive time limits.
41 and 42 · Indicators and costs	OPE indicators; operating and consumption costs.
50 and 51 · People; data and knowledge	Human oversight and capabilities; quality, currency, lineage and retention.
53 · Building solutions with AI	What operations receives at G5.

Version control

Version	Date	Changes
0.1	16-09-2026	First version. Defines the operation of predictive ML (drift, retraining, champion/challenger), generative AI (continuous evaluations, quality, hallucinations, filters, costs, provider versions), agents (actions, limits, intent anomalies, kill switch, permissions) and embedded third-party AI; nine monitoring layers with typical thresholds; three change classes with a return-to- <i>gate</i> table; S1–S4 incident operation with regulatory notifications; R6 package and outcomes; post-market monitoring; continuity and rollback; records and retention; technical retirement; seventeen indicators and the minimum content of P24–P27. Consistency adjustments with 01 (segregation of duties at Lite, R6 outcomes, agents criterion) and with 34 and 37.
0.1	18-09-2026	Usage drift in generative AI and agents (4.2.6); bias in responses with counterfactual pairs (4.2.7); layers 3 and 5 and thresholds extended; activating a planned degraded mode is not a change (6.4); cost-driven degradation cascade NO–N4 with a fallback model validated at G5 (10.3); indicators OPE-18 to OPE-20; Lite and Enterprise requirements and content of P24 and P25.

External sources cited

Links to the official or primary source, checked before publication. Always verify the version in force on the date of use.

Directive (EU) 2022/2555 of the European Parliament and of the Council of 14 December 2022 on measures for a high common level of cybersecurity across the Union (NIS 2 Directive) · <https://eur-lex.europa.eu/eli/dir/2022/2555/oj/eng> · Parlamento Europeo y Consejo de la UE · 14-12-2022

Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation) · <https://eur-lex.europa.eu/eli/reg/2016/679/oj/eng> · Parlamento Europeo y Consejo de la UE · 27-04-2016

Regulation (EU) 2022/2554 of the European Parliament and of the Council of 14 December 2022 on digital operational resilience for the financial sector (DORA) · <https://eur-lex.europa.eu/eli/reg/2022/2554/oj/eng> · Parlamento Europeo y Consejo de la UE · 14-12-2022

Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act) · <https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng> · Parlamento Europeo y Consejo de la UE · 13-06-2024
