

Notes

This presentation answers a very specific question: which data an AI system is built on and how it is kept alive in production.

Many conversations about AI end on go-live day. SEVEN-G starts from the opposite idea: that is where the part that decides whether the value lasts begins. An AI system can fail without going down. It stays available, but its responses get worse, its data changes, its cost grows or its provider changes the model without notice.

The area brings together three moments of the use case lifecycle:

- Data and knowledge, before building: whether they exist, whether their quality is sufficient and whether they can be used lawfully (document 51).
- The build, increasingly assisted by AI: what minimum controls SEVEN-G requires of AI-generated code and how the design gate (G4) and the go-live gate (G5) check them (document 53).
- Operation: monitoring, drift, changes, incidents, cost, rollback and continuity review (document 52).

The ten slides follow this order: why it matters, the key concepts, how SEVEN-G handles it, the governed build, the tools and templates, an example with fictitious figures, common errors, questions for management or the classroom and how to start.

Why it matters: the board approves the investment at G3 and G5, but most of the risk and the value appears afterwards, in daily operation.

Question to open the session: think of an AI system in your organisation that has been in production for a year. Does anyone know whether it works better or worse today than on the day it was approved?

TERMS ON THIS SLIDE

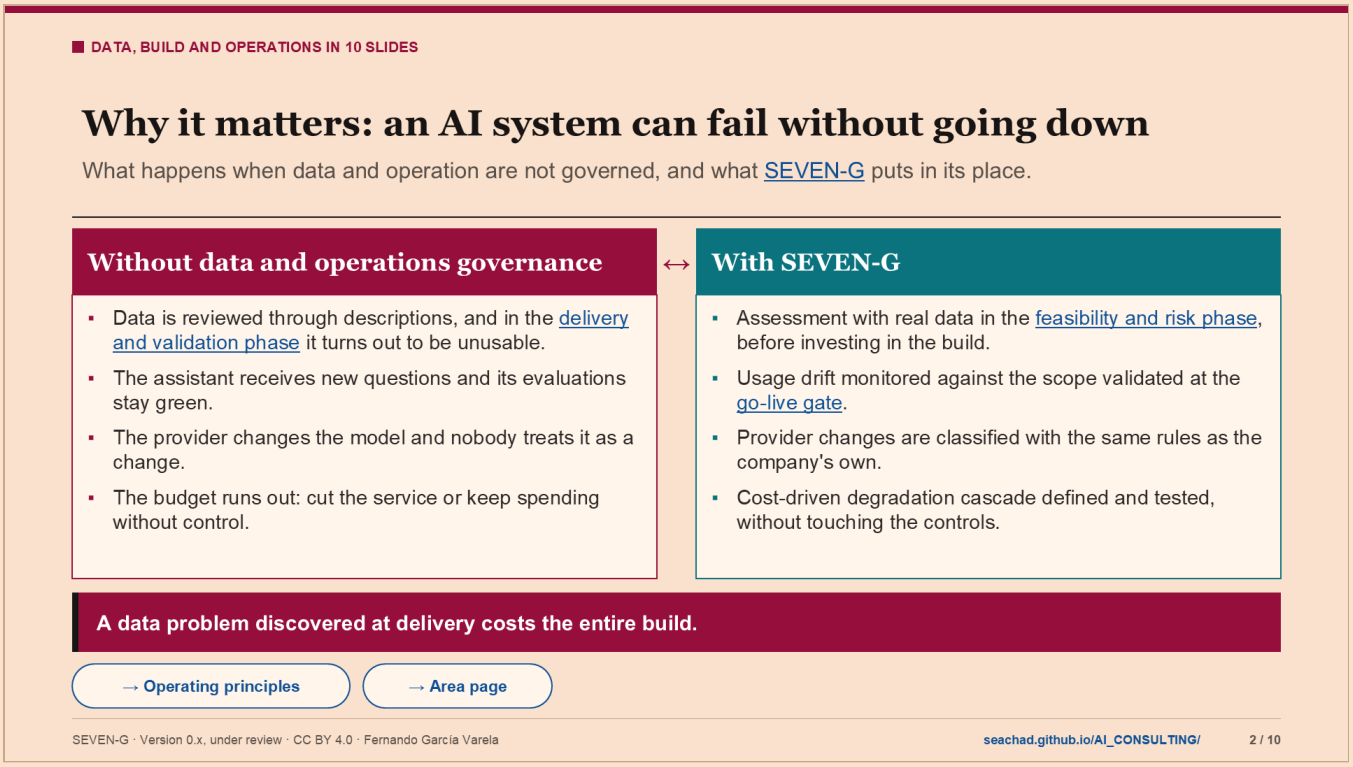
- **AI action areas**: Twelve entry points into the framework by topic: strategy, use cases, value, governance, organisation, maturity, risk, security, standards, board, data and operations, and implementation. Each area answers a question and leads to its documents, templates, tools and a page of its own. It lets you reach a subject

without going through the whole method. The link opens the areas on the SEVEN-G entry page. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/entrada/index.html#areas

- **SEVEN-G**: SEVEN-G (Seven-phase Enterprise Value & Governance) is an open framework for a company to implement, govern and measure AI, and to know whether it is transforming or only gaining efficiency. It treats each use case as an investment with decision gates and gives the board figures it can trust. Management, the AI office, risk and the board use it. The link opens its entry page, with the story, the questions and the examples. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/entrada/index.html

SOURCES AND LINKS

- **Action area: Data, build and operations** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/areas/operacion.html
- **Document 52, section 2: operating principles** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#operating-principles
- **Area page** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/areas/operacion.html



Notes

The left-hand column lists situations that SEVEN-G itself describes as common mistakes. The right-hand column shows what the framework requires to avoid them.

- Data reviewed through descriptions. SEVEN-G starts from a premise: lack of data availability or quality is the main cause of failure of use cases that do not reach production or do not sustain their value. The framework does not base it on market figures, but on its logic, and proposes measuring it in each company with indicator DAT-07. That is why the data assessment is done in phase 3 with real data or a representative sample. An assessment based only on interviews or data dictionaries does not pass the feasibility gate (G3).
- Usage drift. An assistant validated for product queries that starts receiving complaints or legal questions still works and its evaluations stay green, because they measure what was anticipated and not what is happening. SEVEN-G sets a usage reference at G5 and monitors how far use moves away from it.
- The provider changes too. A new version of the model from the same provider and family is a significant change; a change of provider or model family requires a return to G5.
- Budget exhausted. Without a defined cascade, only two bad options remain. With one, cost is contained in steps and no step may disable critical controls.

Why it matters: these failures do not show up in an availability report. They show up in complaints, in bias or in the income statement, when it is already too late.

Question for the classroom: which of the four situations on the left seems most likely in your organisation, and who would detect it?

TERMS ON THIS SLIDE

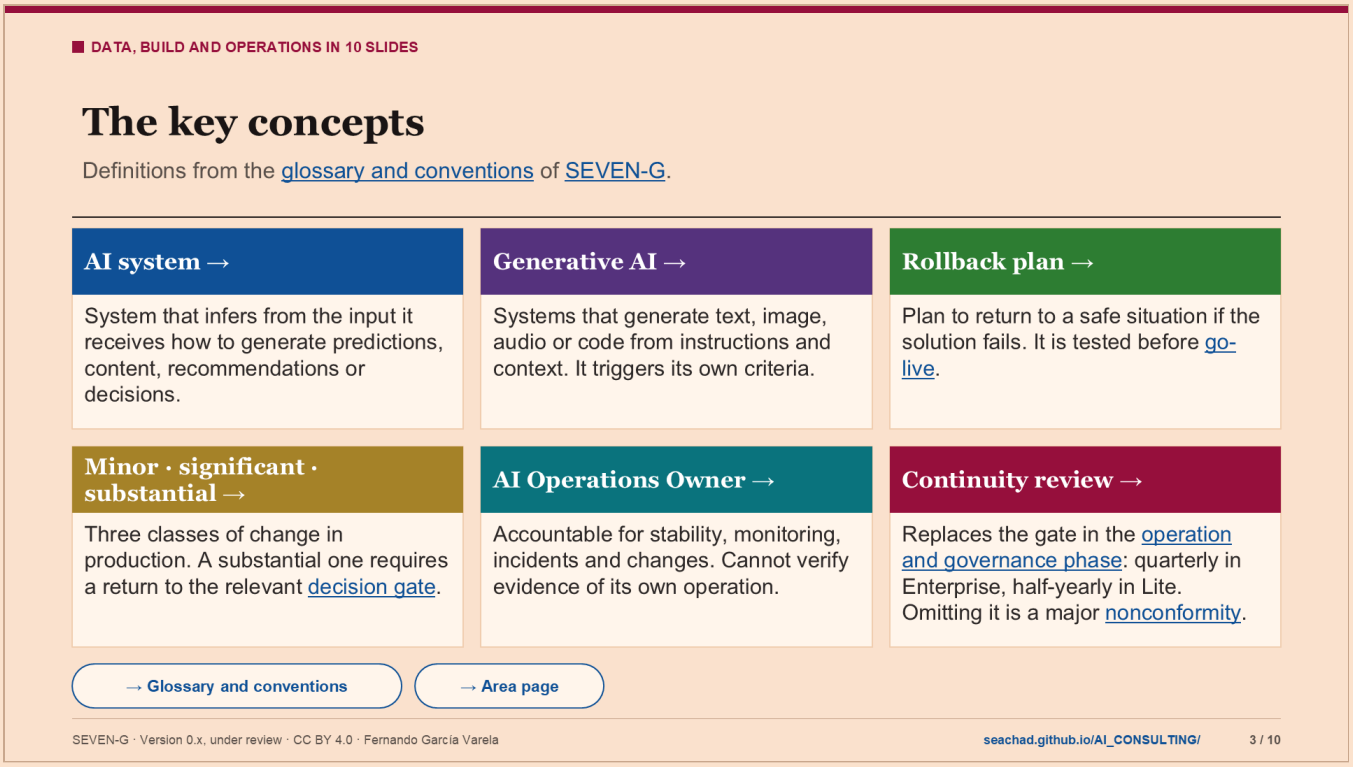
- **SEVEN-G**: SEVEN-G (Seven-phase Enterprise Value & Governance) is an open framework for a company to implement, govern and measure AI, and to know whether it is transforming or only gaining efficiency. It treats each use case as an investment with decision gates and gives the board figures it can trust. Management, the AI

office, risk and the board use it. The link opens its entry page, with the story, the questions and the examples. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/entrada/index.html

- **Delivery and validation** (fase 5): The phase in which the solution is built, tested with real data, piloted and validated to show it delivers what was promised without breaking any constraint. It ends at the go-live gate (G5), with multi-level sign-off. Nothing reaches production without verified evidence. The link opens the phase manual in document 20. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/20_SEVEN-G_Manuales_de_fase.html#phase-5-delivery-and-validation
- **Feasibility and risk** (fase 3): The phase in which a limited test checks whether the use case is technically and economically feasible, it is classified under regulation and its risks are recorded. It ends at the feasibility gate (G3), the main stop gate before the big spend. It is where stopping in time saves the most money. The link opens the phase manual in document 20. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/20_SEVEN-G_Manuales_de_fase.html#phase-3-feasibility-and-risk
- **Go-live gate** (G5): The gate that authorises using the AI system with real users, customers or data. It requires validation and pilot passed, a tested rollback plan and multi-level sign-off (business, technology, risk). Nothing goes into production without this sign-off. The link opens its criteria in document 21. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/21_SEVEN-G_Criterios_de_gate_y_auditoria.html#6-6-g5-go-live

SOURCES AND LINKS

- **Document 51, section 5.2: the phase 3 assessment** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/51_SEVEN-G_Datos_y_conocimiento.html#5-2-phase-3-the-assessment-that-prevents-the-most-stops
- **Document 52, section 4.2.6: usage drift** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#usage-drift
- **Document 52, section 6.2: substantial changes** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#6-2-substantial-changes-and-the-gate-to-return-to
- **Operating principles** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#operating-principles
- **Area page** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/areas/operacion.html



Notes

Six concepts are enough to understand the area.

- AI system: a machine-based system that, with varying levels of autonomy, infers from the input it receives how to generate outputs that can influence physical or virtual environments. The definition is aligned with that of the EU AI Act. Each system has a code in the inventory (SIA-AAAA-NNN).
- Generative AI: systems that generate content from instructions and context. They are operated differently from predictive ML: there are no labels to measure accuracy directly, so they are evaluated with a reference set of cases, responses are sampled and consumption cost is monitored.
- Rollback plan: the plan to return to a safe situation, whether the previous version, a non-AI alternative or reduced autonomy or scope. It is tested before G5 and then at least once a year in Lite and every six months in Enterprise.
- Minor, significant or substantial change: a minor change does not alter behaviour or controls; a significant one alters it within what was approved and requires regression evaluation, updated documentation and a report at the next R6; a substantial one alters the purpose, scope, autonomy, data, critical controls or risk profile and requires a return to the relevant gate.
- AI Operations Owner: the role accountable for stability, monitoring, incidents and changes in production. It cannot verify evidence of its own operation: that is segregation of duties.
- R6 · Continuity review: it replaces the gate in phase 6. It compares realised value with the hypothesis and reviews stability, incidents, compliance and the validity of the classification. Outcomes: Proceed with operation, Proceed with conditions or Bring G7 forward.

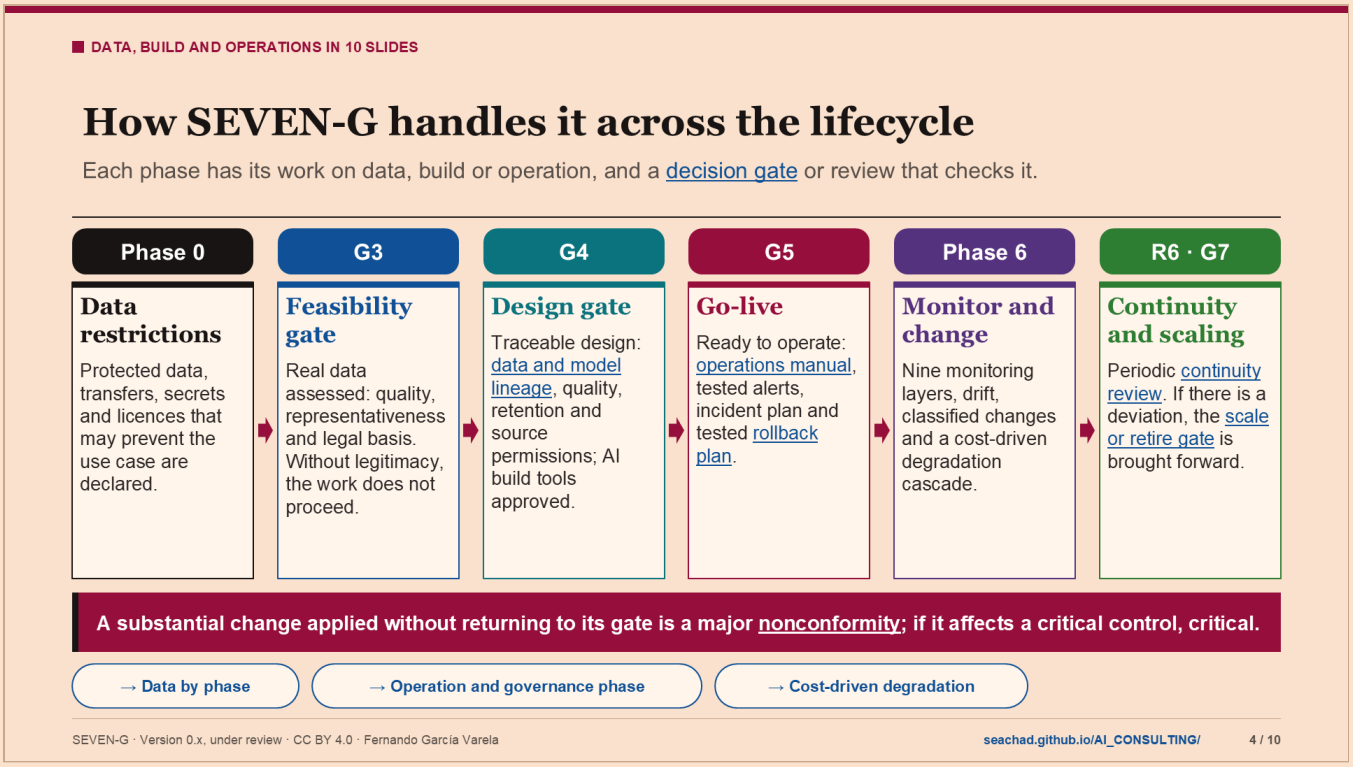
Why it matters: with these six terms, a non-technical management team can ask the essential questions without going into technical detail: which system it is, who is accountable, how to roll back and when it was last reviewed.

TERMS ON THIS SLIDE

- **Glossary and conventions** (documento 02): The official SEVEN-G glossary: each term with its definition, its Spanish equivalent and the document that develops it, plus the scales and codes used across the framework. It ensures everyone calls the same thing by the same name. The link opens the document. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/02_SEVEN-G_Glosario_y_convenciones.html
- **SEVEN-G**: SEVEN-G (Seven-phase Enterprise Value & Governance) is an open framework for a company to implement, govern and measure AI, and to know whether it is transforming or only gaining efficiency. It treats each use case as an investment with decision gates and gives the board figures it can trust. Management, the AI office, risk and the board use it. The link opens its entry page, with the story, the questions and the examples. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/entrada/index.html
- **Go-live gate** (G5): The gate that authorises using the AI system with real users, customers or data. It requires validation and pilot passed, a tested rollback plan and multi-level sign-off (business, technology, risk). Nothing goes into production without this sign-off. The link opens its criteria in document 21. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/21_SEVEN-G_Criterios_de_gate_y_auditoria.html#6-6-g5-go-live
- **Decision gates** (G0–G7): At the end of each phase there is a gate: a moment when a body decides, with verified evidence, whether the use case continues, continues with conditions, iterates, pivots or stops. There are at most two iterations before the decision is escalated. Gates turn enthusiasm into recorded decisions. The link opens their definition in document 01. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/01_SEVEN-G_Metodologia_fundacional.html#decision-gates
- **Operation and governance** (fase 6): The phase in which the use case is in production. Its quality, drift, cost and incidents are monitored, and the value actually obtained is measured. Periodically it goes through a continuity review (R6). It matters because an AI system changes even if nobody touches it. The link opens the phase manual in document 20. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/20_SEVEN-G_Manuales_de_fase.html#phase-6-operation-and-governance
- **Nonconformity** (P50): A breach of a framework rule or a control, even if it caused no harm (for example, a use case in production without sign-off). It is classed as minor, major or critical, recorded (P50), its root cause analysed (P52) and corrected on time. It lets governance learn from its failures. The link opens the process in document 37. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/37_SEVEN-G_No_conformidades_e_incidentes.html#nonconformity-process

SOURCES AND LINKS

- **Document 02: glossary of framework terms** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/02_SEVEN-G_Glosario_y_convenciones.html#glossary-of-framework-terms
- **Document 52, section 6.1: change classes** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#6-1-change-classes
- **Document 52, section 8: R6 continuity review** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#r6-continuity-review
- **Regulation (EU) 2024/1689 (EU AI Act)** https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng
- **Glossary and conventions** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/02_SEVEN-G_Glosario_y_convenciones.html#glossary-of-framework-terms
- **Area page** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/areas/operacion.html



Notes

SEVEN-G does not treat data or operation as separate technical matters: it places them in the lifecycle, with a gate or review that checks them.

- Phase 0. Known data restrictions are declared (P02).
- Phase 3 and the feasibility gate (G3). The data owner, in the business, authorises use for that use case. The assessment covers existence and access, volume, measured quality, labels or ground truth, representativeness, legitimacy, availability in production, knowledge (in generative AI) and preparation cost. Each block is sufficient, sufficient with actions or insufficient.
- Phase 4 and the design gate (G4). Data and model lineage (P16), dataset records, permissions for knowledge sources and AI build tools identified and approved.
- Phase 5 and the go-live gate (G5). Proceed cannot be decided without the operations manual (P24), monitoring with tested alerts (P25), the incident plan (P26), the change and incident log (P27) and the tested rollback plan (P19). The operations owner takes part in phase 5 and accepts the handover.
- Phase 6. The operations owner runs the system. Nine layers are monitored. Drift takes three forms in predictive ML (data, prediction and concept) and one more in generative AI and agents: usage drift. Each change is classified; significant ones in Enterprise need risk clearance from the AI Risk Owner. If the consumption budget runs out, the cascade moves from N0 (normal) to N1 (optimise), N2 (fallback model), N3 (reduced scope) and N4 (alternative without AI). The model is not downgraded in uses that decide about people without specific validation.
- R6 and G7. In Lite the AI Office and the sponsor decide; in Enterprise, the AI Auditor and the AI Committee.

Why it matters: each control arrives at the moment when it is still cheap.

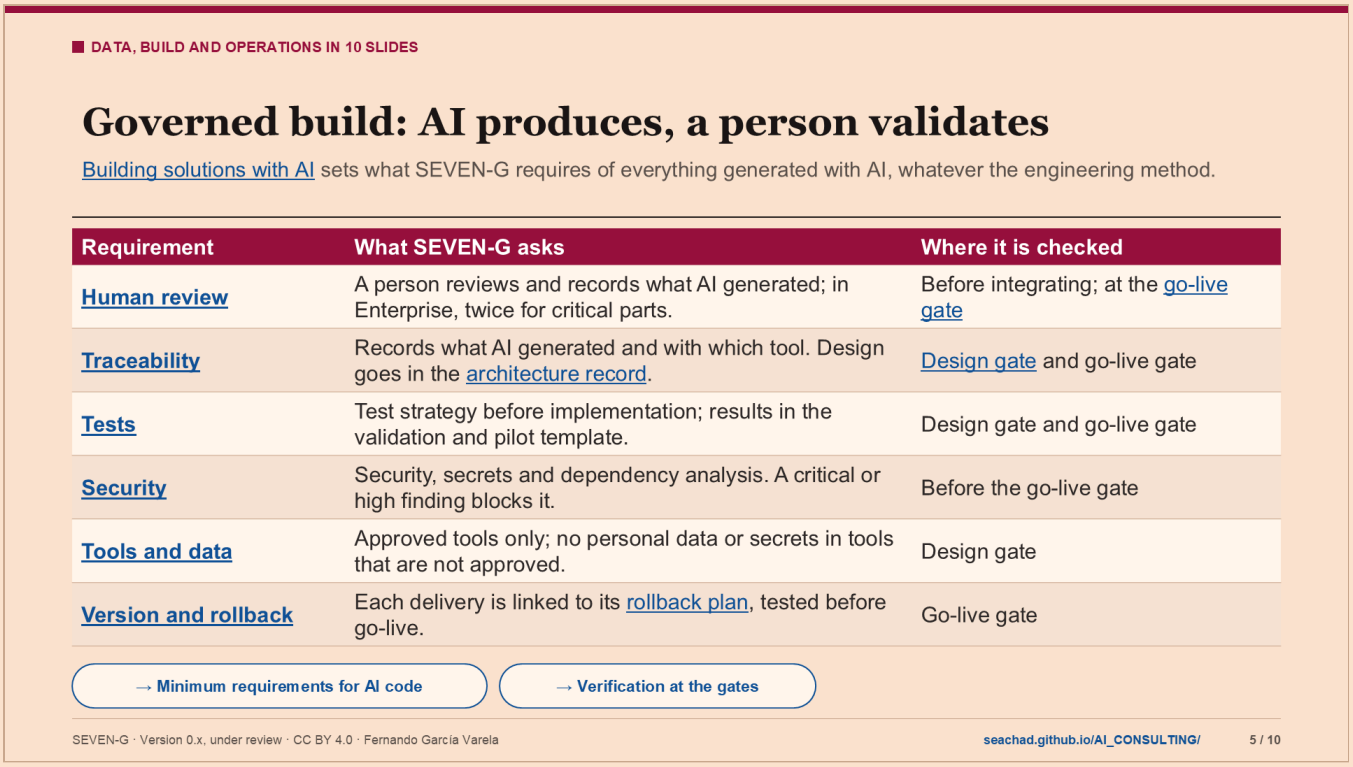
Question for the classroom: in which phase is it usually discovered today that the data is not fit for use?

TERMS ON THIS SLIDE

- **Decision gates** (G0–G7): At the end of each phase there is a gate: a moment when a body decides, with verified evidence, whether the use case continues, continues with conditions, iterates, pivots or stops. There are at most two iterations before the decision is escalated. Gates turn enthusiasm into recorded decisions. The link opens their definition in document 01. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/01_SEVEN-G_Metodologia_fundacional.html#decision-gates
- **Data and model lineage** (P16): A template that records where the data comes from, which models and versions are used, how they were trained or configured and which fallback model exists. It makes it possible to reconstruct what happened if something fails and to change model without surprises. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P16_SEVEN-G_Linaje_de_datos_y_modelos.html
- **Operations manual** (P24): A template with everything needed to run the AI system in production: owners, daily tasks, thresholds, what to do on an alert, how to escalate and how to trigger rollback. It is evidence for go-live. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P24_SEVEN-G_Manual_de_operacion.html
- **Rollback plan** (P19): A tested plan to go back if the AI system fails: how it is switched off, which non-AI alternative is used and how fast. For agents it includes a kill switch. It is designed before go-live and actually tested (template P19). The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P19_SEVEN-G_Plan_de_reversion_y_prueba.html
- **Continuity review** (R6): A periodic review of a use case in production: quarterly in Enterprise, half-yearly in Lite and yearly in Express. It checks whether value is still arriving and whether there is drift, incidents or regulatory change. A deviation brings forward the scale or retire decision (G7). It is prepared with pack P65. The link opens its criteria in document 21. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/21_SEVEN-G_Criterios_de_gate_y_auditoria.html#6-7-r6-continuity-review
- **Scale or retire gate** (G7): The final gate of the lifecycle. With the value actually measured, it decides to scale the use case (which starts a new cycle), iterate it or retire it with an orderly plan. Scaling a Transform bet is approved by the board. The link opens its criteria in document 21. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/21_SEVEN-G_Criterios_de_gate_y_auditoria.html#6-8-g7-scale-or-retire
- **Nonconformity** (P50): A breach of a framework rule or a control, even if it caused no harm (for example, a use case in production without sign-off). It is classed as minor, major or critical, recorded (P50), its root cause analysed (P52) and corrected on time. It lets governance learn from its failures. The link opens the process in document 37. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/37_SEVEN-G_No_conformidades_e_incidentes.html#nonconformity-process

SOURCES AND LINKS

- **Document 51, section 5: data preparation by phase** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/51_SEVEN-G_Datos_y_conocimiento.html#data-preparation-by-phase
- **Document 52, section 3.2: readiness to operate** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#3-2-readiness-to-operate
- **Document 52, section 10.3: cost-driven degradation cascade** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#10-3-cost-driven-degradation-cascade
- **Document 20, section 9: phase 6, operation and governance** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/20_SEVEN-G_Manuales_de_fase.html#phase-6-operation-and-governance
- **Data by phase** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/51_SEVEN-G_Datos_y_conocimiento.html#5-1-requirements-by-phase
- **Operation and governance phase** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/20_SEVEN-G_Manuales_de_fase.html#phase-6-operation-and-governance
- **Cost-driven degradation** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#10-3-cost-driven-degradation-cascade



DATA, BUILD AND OPERATIONS IN 10 SLIDES

Governed build: AI produces, a person validates

Building solutions with AI sets what SEVEN-G requires of everything generated with AI, whatever the engineering method.

Requirement	What SEVEN-G asks	Where it is checked
Human review	A person reviews and records what AI generated; in Enterprise, twice for critical parts.	Before integrating; at the go-live gate
Traceability	Records what AI generated and with which tool. Design goes in the architecture record.	Design gate and go-live gate
Tests	Test strategy before implementation; results in the validation and pilot template.	Design gate and go-live gate
Security	Security, secrets and dependency analysis. A critical or high finding blocks it.	Before the go-live gate
Tools and data	Approved tools only; no personal data or secrets in tools that are not approved.	Design gate
Version and rollback	Each delivery is linked to its rollback plan, tested before go-live.	Go-live gate

→ Minimum requirements for AI code → Verification at the gates

SEVEN-G · Version 0.x, under review · CC BY 4.0 · Fernando García Varela seachad.github.io/AI_CONSULTING/ 5 / 10

Notes

More and more solutions are built with coding assistants and agents that write code. This happens whether or not the solution is an AI system. SEVEN-G separates two objects: the system being built, which is governed through the lifecycle, and the AI tools that help build it, which are inventoried and approved like any other use of AI.

Document 53 sets the minimum requirements that SEVEN-G applies to all code, configuration, tests or technical documentation generated with AI, whatever the engineering method the team uses. They apply mainly in phases 4 (design) and 5 (delivery):

- Human review. A competent person reviews every AI-generated artefact before it is integrated and records who, when and with what result. In Enterprise, critical components such as payments, personal data or agent limits get a second review. An automated review does not replace the human one. No AI approves.
- Traceability. It is known which part was generated with AI and with which tool, and design decisions are recorded in P15, not only in conversations with the tool.
- Tests. The strategy is defined before implementation, and tests generated with AI are reviewed to check that they verify the expected behaviour.
- Security. Static analysis, secret detection and dependency review before G5. Critical and high findings block G5.
- Tools and data. Approved tools only, and no personal data or secrets in a tool that is not approved.
- Version and rollback. Each delivery is linked to its rollback plan (P19).

Who validates: the design gate (G4) checks that the tools are approved, that the plan and the rules are recorded in P15 and that the test strategy exists. The go-live gate (G5) checks the recorded human review, the test and security results in P22 with no critical or high findings open, the licences and the tested rollback. The technical annexes are kept in the build annexes template (P66).

Why it matters: AI generates plausible code at great speed. Without these requirements, that speed turns into design decisions that nobody took consciously.

TERMS ON THIS SLIDE

- **Building solutions with AI** (documento 53): What SEVEN-G requires of how a use case's solution is built: testing, code review, evaluation of model behaviour, change control and the evidence handed to the go-live gate. It refers to a separate building methodology. The link opens the document. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/53_SEVEN-G_Construccion_de_soluciones_con_IA.html
- **Go-live gate** (G5): The gate that authorises using the AI system with real users, customers or data. It requires validation and pilot passed, a tested rollback plan and multi-level sign-off (business, technology, risk). Nothing goes into production without this sign-off. The link opens its criteria in document 21. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/21_SEVEN-G_Criterios_de_gate_y_auditoria.html#6-6-g5-go-live
- **Architecture record** (P15): A template describing the solution architecture: components, models, integrations, data flows and technical decisions with their reasons. It is filled in during the design phase. It allows risks and security to be reviewed against something concrete and lets another team maintain it. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P15_SEVEN-G_Registro_de_arquitectura.html
- **Design gate** (G4): The gate that approves the full design before building: architecture, data, security, human oversight, rollback and adoption. It checks that the controls are planned from the start. The link opens its criteria in document 21. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/21_SEVEN-G_Criterios_de_gate_y_auditoria.html#6-5-g4-design
- **Rollback plan** (P19): A tested plan to go back if the AI system fails: how it is switched off, which non-AI alternative is used and how fast. For agents it includes a kill switch. It is designed before go-live and actually tested (template P19). The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P19_SEVEN-G_Plan_de_reversion_y_prueba.html

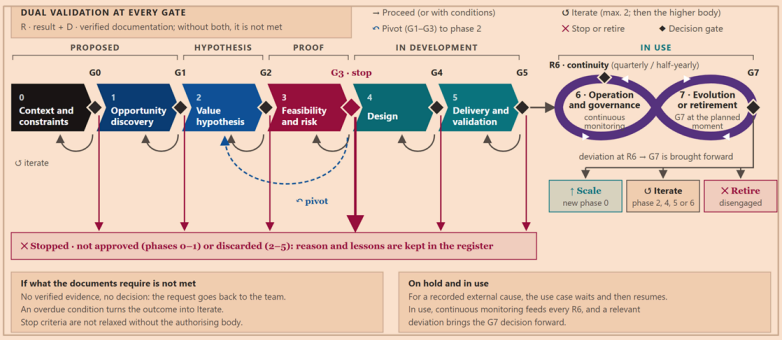
SOURCES AND LINKS

- **Document 53, section 7: minimum requirements for AI-generated code** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/53_SEVEN-G_Construccion_de_soluciones_con_IA.html#seven-g-minimum-requirements-for-ai-generated-code
- **Document 53, section 8: verification at the gates** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/53_SEVEN-G_Construccion_de_soluciones_con_IA.html#verification-at-the-gates
- **Document 53, section 1.3: two distinct objects** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/53_SEVEN-G_Construccion_de_soluciones_con_IA.html#1-3-two-distinct-objects
- **Minimum requirements for AI code** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/53_SEVEN-G_Construccion_de_soluciones_con_IA.html#seven-g-minimum-requirements-for-ai-generated-code
- **Verification at the gates** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/53_SEVEN-G_Construccion_de_soluciones_con_IA.html#verification-at-the-gates

DATA, BUILD AND OPERATIONS IN 10 SLIDES

Tools and templates used

All are adaptable. The figure shows the operation loop: monitoring, continuity review and scale or retire.



- **Data and lineage** The dataset record and knowledge source inventory and the data and model lineage.
- **Operations and alerts** The operations manual and the monitoring and alerting configuration, with thresholds, owners and tested alerts.
- **Rollback and log** The rollback plan and its test; the incident and change log with the class of each change.
- **Continuity** Continuity review package; the use case register warns when a review is overdue.

→ Operations templates | → Use case register

SEVEN-G · Version 0.x, under review · CC BY 4.0 · Fernando García Varela seachad.github.io/AI_CONSULTING/ 6 / 10

Notes

The area does not create new tools: it uses templates that the company adapts and, if it already has a data catalogue or a change management system, relies on them.

- Dataset record and inventory of knowledge sources (P64). Every dataset used to train, validate or feed an AI system has an owner, origin, legal basis, measured quality, representativeness, retention and version. Generative AI knowledge sources have a content owner, permissions and a review date.
- Data and model lineage (P16). It is prepared in phase 4, must be up to date for G5 and is updated with every change.
- Operations manual (P24). It describes the system, its autonomy, its service levels, the on-call arrangements, the degraded modes and the cost-driven degradation cascade with its triggers.
- Monitoring and alert configuration (P25). For each layer, the metric, the baseline, the warning and critical thresholds, the initial severity and the owner. Alerts are tested before G5.
- Rollback plan and test (P19). How to return to the previous version or to a non-AI alternative, and the result of the test.
- Incident and change log (P27). Each change with its class and approval; each incident with its severity S1 to S4, its root cause and its closure.
- Continuity review package (P65). Value against the hypothesis, cost, stability, drift, incidents, agents, human oversight, risk and continuity, with a proposal.
- Use case register (T01). It holds the phase and status of each use case; an overdue R6 generates an alert.

The figure shows how a use case flows: after go-live it enters an operation loop with continuous monitoring and a periodic R6, which can bring forward evolution or retirement (G7).

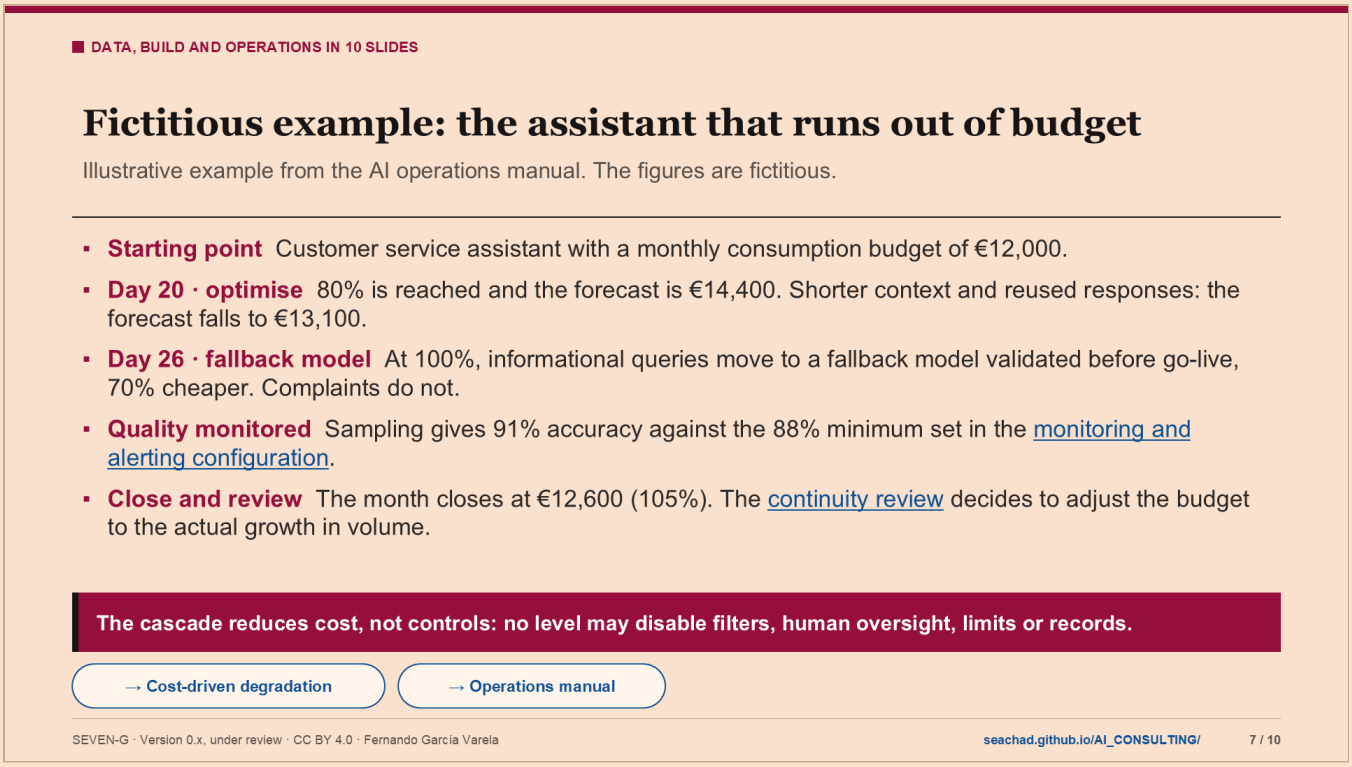
Why it matters: each template leaves the evidence that an auditor or a regulator will ask for. Without it, what is not recorded did not happen.

TERMS ON THIS SLIDE

- **Dataset record and knowledge source inventory** (P64): A template to describe each dataset and each knowledge source AI uses: origin, owner, quality, permissions, personal data and review date. Assistants that query internal documents depend on these sources being up to date. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P64_SEVEN-G_Ficha_de_conjunto_de_datos_e_inventario_de_fuentes_de_conocimiento.html
- **Data and model lineage** (P16): A template that records where the data comes from, which models and versions are used, how they were trained or configured and which fallback model exists. It makes it possible to reconstruct what happened if something fails and to change model without surprises. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P16_SEVEN-G_Linaje_de_datos_y_modelos.html
- **Operations manual** (P24): A template with everything needed to run the AI system in production: owners, daily tasks, thresholds, what to do on an alert, how to escalate and how to trigger rollback. It is evidence for go-live. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P24_SEVEN-G_Manual_de_operacion.html
- **Monitoring and alerting configuration** (P25): A template that sets what is monitored in production (quality, data and usage drift, bias, cost, control effectiveness) and at which thresholds an alert fires and to whom it goes. An AI system changes even if nobody touches it. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P25_SEVEN-G_Configuracion_de_monitorizacion_y_alertas.html
- **Rollback plan** (P19): A tested plan to go back if the AI system fails: how it is switched off, which non-AI alternative is used and how fast. For agents it includes a kill switch. It is designed before go-live and actually tested (template P19). The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P19_SEVEN-G_Plan_de_reversion_y_prueba.html
- **Incident and change log** (P27): A template that logs the incidents and changes of each system in production: what happened or what changed, when, who approved it and what was learned. It is the system's memory and a key input to the continuity review. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P27_SEVEN-G_Registro_de_incidentes_y_cambios.html
- **Use case register** (T01): It is the central SEVEN-G tool and the single data source for all the others. It records each AI use case like an opportunity in a CRM: phase, status, dates, criteria for each gate, risks, value and stop reason. The AI office and the teams use it; the committee and the board see what comes out of it. The link opens an example register with fictitious use cases; anything you change is stored only in your browser. https://seachad.github.io/AI_CONSULTING/SEVEN-G/herramientas/T01_registro_iniciativas/registro.html

SOURCES AND LINKS

- **Template P65: R6 continuity review package** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P65_SEVEN-G_Paquete_de_revision_de_continuidad_R6.html
- **Operations templates** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#15-1-operations-templates



Notes

This example is fictitious and is taken from SEVEN-G document 52 (section 10.3). It shows how the cost of a generative AI system is contained without improvisation.

Systems with variable consumption, such as generative AI and agents, must have defined in their operations manual (P24) what happens when consumption reaches the budget. The response follows a cascade of levels, from least to most restrictive: N0 normal operation; N1 optimisation without a model change; N2 fallback model, evaluated before G5 with the same evaluation set, including the bias tests; N3 reduced scope, in which people handle the rest; and N4 alternative without AI.

What the example shows:

- The response was decided beforehand. The triggers (N1 at 80%, N2 at 100%) and who activates them are set in P24 and were tested before G5. Activating a planned degraded mode is not a change: it is an approved operating mode and is recorded as an event in P27.
- Not everything is degraded in the same way. Informational queries move to the fallback model; complaints stay on the primary model. In uses that decide or recommend about people, in critical functions or in high-risk systems, the model is not downgraded without specific validation: the system moves to N3 or N4, or a budget supplement is approved.
- Quality is monitored more closely during the mode: sampling is at least doubled and, if quality falls below the minimum, the system moves to the next level.
- Learning is part of the process. If a system spends more than a month in N2 or a higher level, the budget or the design is wrongly sized and it is taken to the R6 with a proposal.

Why it matters: the board knows in advance what quality is delivered at each level and what each month costs, instead of finding out through an invoice or a service outage.

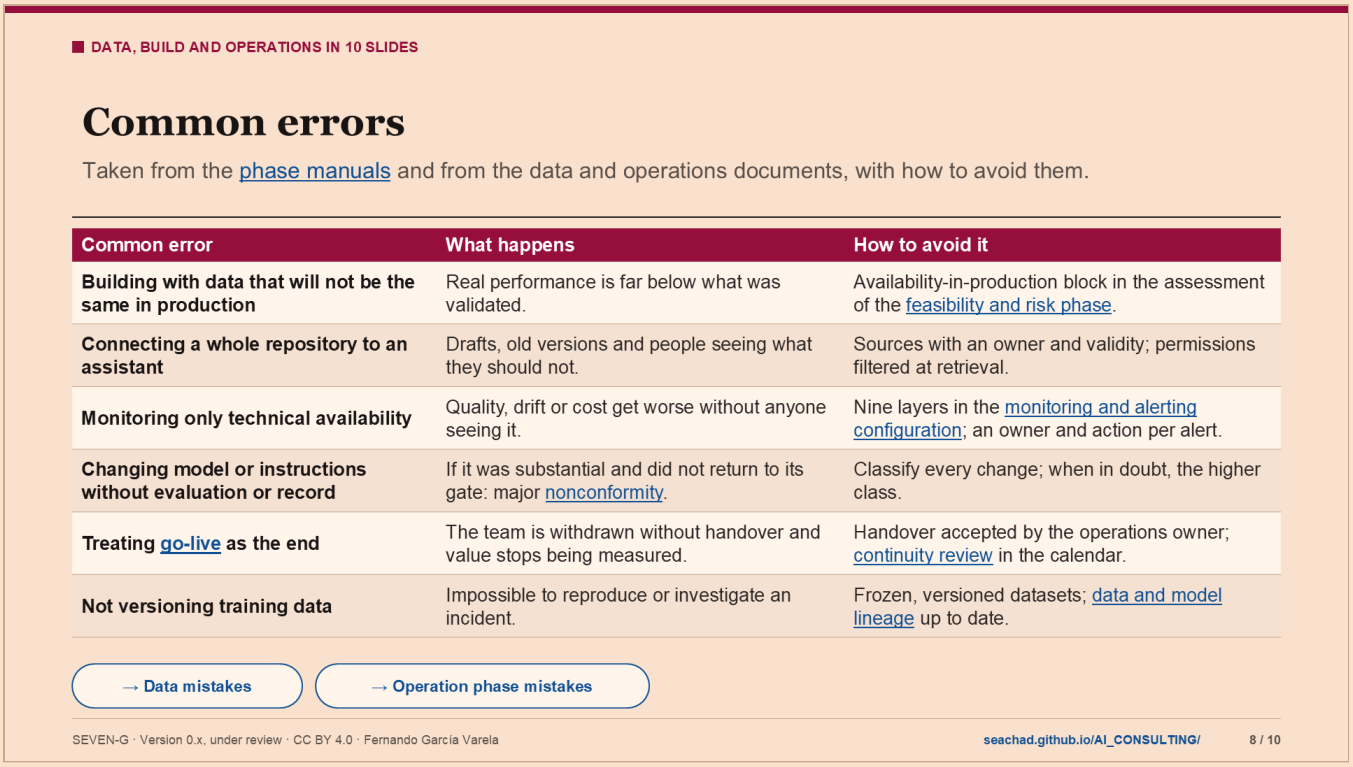
Question for the classroom: which use cases in your organisation should never move to a cheaper model, and why?

TERMS ON THIS SLIDE

- **Monitoring and alerting configuration** (P25): A template that sets what is monitored in production (quality, data and usage drift, bias, cost, control effectiveness) and at which thresholds an alert fires and to whom it goes. An AI system changes even if nobody touches it. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P25_SEVEN-G_Configuracion_de_monitorizacion_y_alertas.html
- **Continuity review** (R6): A periodic review of a use case in production: quarterly in Enterprise, half-yearly in Lite and yearly in Express. It checks whether value is still arriving and whether there is drift, incidents or regulatory change. A deviation brings forward the scale or retire decision (G7). It is prepared with pack P65. The link opens its criteria in document 21. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/21_SEVEN-G_Criterios_de_gate_y_auditoria.html#6-7-r6-continuity-review

SOURCES AND LINKS

- **Document 52, section 10.3: cost-driven degradation cascade** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#10-3-cost-driven-degradation-cascade
- **Document 52, section 6.4: change rules** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#6-4-rules
- **Cost-driven degradation** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#10-3-cost-driven-degradation-cascade
- **Operations manual** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P24_SEVEN-G_Manual_de_operacion.html



DATA, BUILD AND OPERATIONS IN 10 SLIDES

Common errors

Taken from the phase manuals and from the data and operations documents, with how to avoid them.

Common error	What happens	How to avoid it
Building with data that will not be the same in production	Real performance is far below what was validated.	Availability-in-production block in the assessment of the feasibility and risk phase.
Connecting a whole repository to an assistant	Drafts, old versions and people seeing what they should not.	Sources with an owner and validity; permissions filtered at retrieval.
Monitoring only technical availability	Quality, drift or cost get worse without anyone seeing it.	Nine layers in the monitoring and alerting configuration; an owner and action per alert.
Changing model or instructions without evaluation or record	If it was substantial and did not return to its gate: major nonconformity.	Classify every change; when in doubt, the higher class.
Treating go-live as the end	The team is withdrawn without handover and value stops being measured.	Handover accepted by the operations owner; continuity review in the calendar.
Not versioning training data	Impossible to reproduce or investigate an incident.	Frozen, versioned datasets; data and model lineage up to date.

→ Data mistakes → Operation phase mistakes

SEVEN-G · Version 0.x, under review · CC BY 4.0 · Fernando García Varela seachad.github.io/AI_CONSULTING/ 8 / 10

Notes

The six errors in the table come from the common mistakes lists in document 51 (data and knowledge) and in the phase 6 manual (document 20), and from the change rules in document 52.

- Data that will not be the same in production. A model is built on a clean historical extract and in production receives data with another definition, another frequency or another delay. That is why the phase 3 assessment includes a specific block on availability in production.
- A whole repository connected to an assistant. The quality of generative AI responses depends more on the sources than on the model. A source should not be connected if it has no content owner, if it mixes drafts and official documents or if its permissions cannot be carried over. Permissions are filtered at the moment of retrieval, not only at indexing. If the assistant cites an expired document on prices or terms, it is an incident.
- Availability only. The first principle of the operations manual says it: an AI system can fail without going down.
- Changes without evaluation. Retraining, changing the system instructions or adopting a new version of the provider's model are significant changes. Increasing autonomy, for example from recommendation to supervised action, is substantial and requires a return to G4 and G5.
- Go-live as the end. Phase 6 has no fixed end. Omitting an R6 is a major nonconformity.
- Unversioned data. Without frozen datasets, a result cannot be reproduced and a regulator cannot be answered.

Why it matters: none of these errors is technical in origin. They are organisational decisions that someone did not take in time.

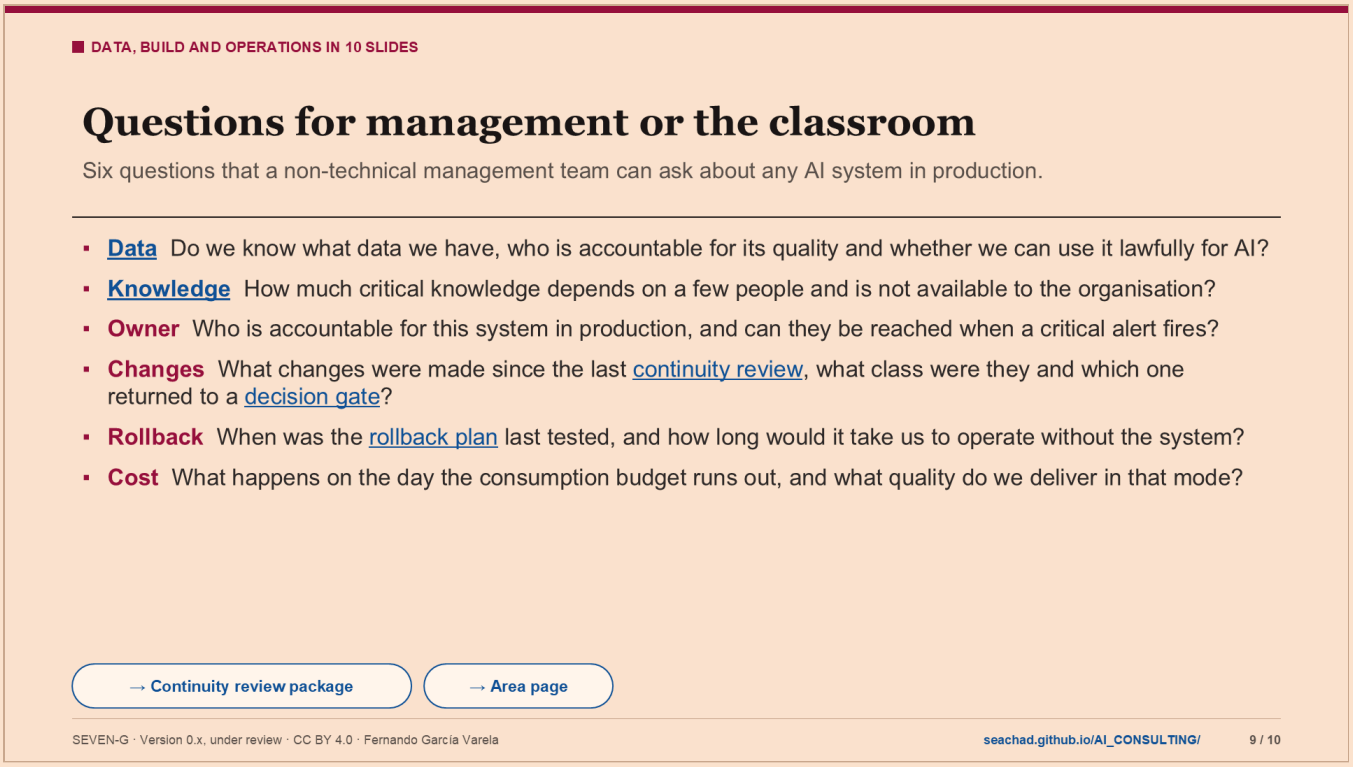
Question for the classroom: which of these errors would be most expensive to correct in your organisation if it were discovered a year from now?

TERMS ON THIS SLIDE

- **Phase manuals** (documento 20): One manual for each lifecycle phase (0 to 7): objective, activities, who takes part, which templates are filled in and which gate closes it, plus what happens if it is not passed. It is the working guide for each use case team. The link opens the document. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/20_SEVEN-G_Manuales_de_fase.html
- **Feasibility and risk** (fase 3): The phase in which a limited test checks whether the use case is technically and economically feasible, it is classified under regulation and its risks are recorded. It ends at the feasibility gate (G3), the main stop gate before the big spend. It is where stopping in time saves the most money. The link opens the phase manual in document 20. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/20_SEVEN-G_Manuales_de_fase.html#phase-3-feasibility-and-risk
- **Monitoring and alerting configuration** (P25): A template that sets what is monitored in production (quality, data and usage drift, bias, cost, control effectiveness) and at which thresholds an alert fires and to whom it goes. An AI system changes even if nobody touches it. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P25_SEVEN-G_Configuracion_de_monitorizacion_y_alertas.html
- **Nonconformity** (P50): A breach of a framework rule or a control, even if it caused no harm (for example, a use case in production without sign-off). It is classed as minor, major or critical, recorded (P50), its root cause analysed (P52) and corrected on time. It lets governance learn from its failures. The link opens the process in document 37. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/37_SEVEN-G_No_conformidades_e_incidentes.html#nonconformity-process
- **Go-live gate** (G5): The gate that authorises using the AI system with real users, customers or data. It requires validation and pilot passed, a tested rollback plan and multi-level sign-off (business, technology, risk). Nothing goes into production without this sign-off. The link opens its criteria in document 21. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/21_SEVEN-G_Criterios_de_gate_y_auditoria.html#6-6-g5-go-live
- **Continuity review** (R6): A periodic review of a use case in production: quarterly in Enterprise, half-yearly in Lite and yearly in Express. It checks whether value is still arriving and whether there is drift, incidents or regulatory change. A deviation brings forward the scale or retire decision (G7). It is prepared with pack P65. The link opens its criteria in document 21. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/21_SEVEN-G_Criterios_de_gate_y_auditoria.html#6-7-r6-continuity-review
- **Data and model lineage** (P16): A template that records where the data comes from, which models and versions are used, how they were trained or configured and which fallback model exists. It makes it possible to reconstruct what happened if something fails and to change model without surprises. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P16_SEVEN-G_Linaje_de_datos_y_modelos.html

SOURCES AND LINKS

- **Document 51, section 13: common mistakes** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/51_SEVEN-G_Datos_y_conocimiento.html#common-mistakes
- **Document 51, section 10.4: access permissions** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/51_SEVEN-G_Datos_y_conocimiento.html#10-4-access-permissions
- **Document 20, section 9.12: phase 6 common mistakes** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/20_SEVEN-G_Manuales_de_fase.html#9-12-common-mistakes
- **Document 52, section 6.4: change rules** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#6-4-rules
- **Data mistakes** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/51_SEVEN-G_Datos_y_conocimiento.html#common-mistakes
- **Operation phase mistakes** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/20_SEVEN-G_Manuales_de_fase.html#9-12-common-mistakes



Notes

The first two questions are those that SEVEN-G proposes to the board for enabling spheres 05 (Data) and 06 (Knowledge) in document 51. The other four come from the operations manual (document 52) and from what the continuity review (R6) package must contain.

How to use them:

- In a management or board session, ask for the answer to come with evidence: the dataset record (P64), the operations manual (P24), the change log (P27), the result of the last rollback test (P19) or the definition of the cost-driven degradation cascade.
- An answer such as “the technical team handles it” is not enough. SEVEN-G requires an operations owner who can be reached in accordance with the service level; no owner, no production.
- The question on critical knowledge has a people dimension. An expert's knowledge is captured with their voluntary and informed participation, with recognition and with clarity about later use. It must not be presented as a step before letting the person go without assessing the effect in line with document 50.
- The question on rollback has two parts: the technical test and the people. The plan identifies who, and with what capabilities, operates without the system.

In the classroom, these questions help to analyse a case: the group chooses a well-known AI system, answers the six questions with what it knows and points out which evidence it would be missing.

Why it matters: good questions do not require technical knowledge. They require knowing which evidence should exist and asking for it.

TERMS ON THIS SLIDE

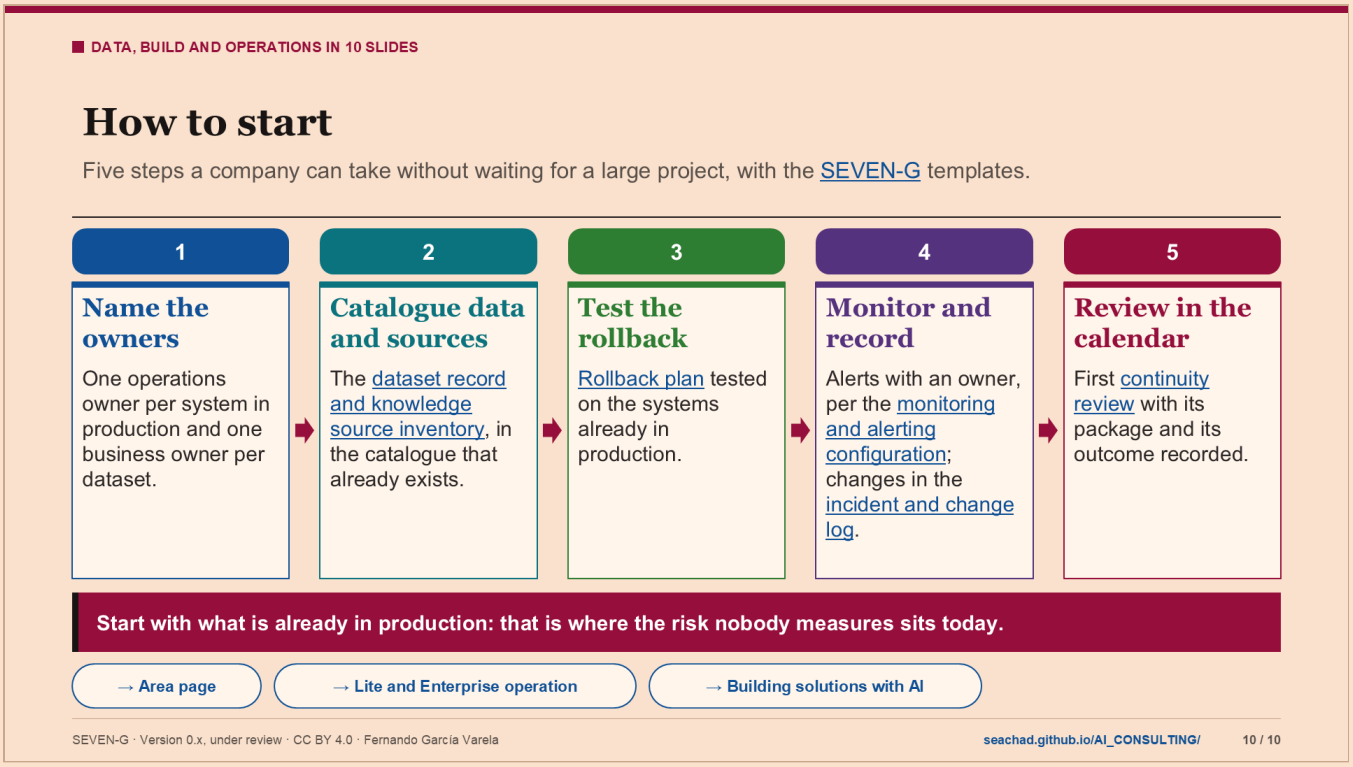
- **Continuity review** (R6): A periodic review of a use case in production: quarterly in Enterprise, half-yearly in Lite and yearly in Express. It checks whether value is still arriving and whether there is drift, incidents or regulatory change. A deviation brings forward the scale or retire decision (G7). It is prepared with pack P65. The link opens its

criteria in document 21. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/21_SEVEN-G_Criterios_de_gate_y_auditoria.html#6-7-r6-continuity-review

- **Decision gates** (G0–G7): At the end of each phase there is a gate: a moment when a body decides, with verified evidence, whether the use case continues, continues with conditions, iterates, pivots or stops. There are at most two iterations before the decision is escalated. Gates turn enthusiasm into recorded decisions. The link opens their definition in document 01. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/01_SEVEN-G_Metodologia_fundacional.html#decision-gates
- **Rollback plan** (P19): A tested plan to go back if the AI system fails: how it is switched off, which non-AI alternative is used and how fast. For agents it includes a kill switch. It is designed before go-live and actually tested (template P19). The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P19_SEVEN-G_Plan_de_reversion_y_prueba.html

SOURCES AND LINKS

- **Document 51, section 2.1: reference questions** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/51_SEVEN-G_Datos_y_conocimiento.html#2-1-reference-questions
- **Document 52, section 8.2: R6 package** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#8-2-r6-package
- **Document 52, section 10.2: continuity and rollback rules** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#10-2-rules
- **Continuity review package** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#8-2-r6-package
- **Area page** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/areas/operacion.html



Notes

These five steps order what SEVEN-G already requires in documents 51 and 52, so that a company can start with what is most urgent. They do not create new rules.

- Step 1. Name the owners. Two principles of the framework require it: no data without an owner, and no owner, no production. The operations owner must be reachable during the system's operating hours.
- Step 2. Catalogue data and sources. If the company already has a data catalogue, the record is implemented as an extension of that catalogue; if not, with template P64. For generative AI, each knowledge source needs a content owner, permissions and a review date.
- Step 3. Test the rollback. At least once a year in Lite and every six months in Enterprise, and after every substantial change. The plan also identifies the people who would operate without the system.
- Step 4. Monitor and record. In Lite a reduced set of layers is enough: availability, input data, quality or performance depending on the type of system, cost, and human oversight with value. In Enterprise, all those that apply. Every alert has an owner, a procedure and an initial severity.
- Step 5. Put R6 in the calendar: half-yearly in Lite and quarterly in Enterprise. Its possible outcomes are Proceed with operation, Proceed with conditions or Bring G7 forward.

For systems that were already in production before the framework was applied, SEVEN-G provides for their regularisation (document 20, section 11.8). For what is built from now on with AI assistance, the company applies the minimum requirements of document 53, with the engineering method it chooses.

All the material is free and adaptable (content CC BY 4.0 and code MIT). SEVEN-G does not certify: the company applies the framework and verifies its own compliance.

Why it matters: AI risk does not wait for the governance project to finish. The systems already in operation are the first ones that must be explainable.

TERMS ON THIS SLIDE

- **SEVEN-G**: SEVEN-G (Seven-phase Enterprise Value & Governance) is an open framework for a company to implement, govern and measure AI, and to know whether it is transforming or only gaining efficiency. It treats each use case as an investment with decision gates and gives the board figures it can trust. Management, the AI office, risk and the board use it. The link opens its entry page, with the story, the questions and the examples. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/entrada/index.html
- **Dataset record and knowledge source inventory** (P64): A template to describe each dataset and each knowledge source AI uses: origin, owner, quality, permissions, personal data and review date. Assistants that query internal documents depend on these sources being up to date. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P64_SEVEN-G_Ficha_de_conjunto_de_datos_e_inventario_de_fuentes_de_conocimiento.html
- **Rollback plan** (P19): A tested plan to go back if the AI system fails: how it is switched off, which non-AI alternative is used and how fast. For agents it includes a kill switch. It is designed before go-live and actually tested (template P19). The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P19_SEVEN-G_Plan_de_reversion_y_prueba.html
- **Monitoring and alerting configuration** (P25): A template that sets what is monitored in production (quality, data and usage drift, bias, cost, control effectiveness) and at which thresholds an alert fires and to whom it goes. An AI system changes even if nobody touches it. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P25_SEVEN-G_Configuracion_de_monitorizacion_y_alertas.html
- **Incident and change log** (P27): A template that logs the incidents and changes of each system in production: what happened or what changed, when, who approved it and what was learned. It is the system's memory and a key input to the continuity review. The link opens the template. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P27_SEVEN-G_Registro_de_incidentes_y_cambios.html
- **Continuity review** (R6): A periodic review of a use case in production: quarterly in Enterprise, half-yearly in Lite and yearly in Express. It checks whether value is still arriving and whether there is drift, incidents or regulatory change. A deviation brings forward the scale or retire decision (G7). It is prepared with pack P65. The link opens its criteria in document 21. https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/21_SEVEN-G_Criterios_de_gate_y_auditoria.html#6-7-r6-continuity-review

SOURCES AND LINKS

- **Document 52, section 14: Lite and Enterprise operation** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#lite-and-enterprise-operation
- **Document 51, section 3: principles** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/51_SEVEN-G_Datos_y_conocimiento.html#principles
- **Template P65: R6 continuity review package** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/plantillas/P65_SEVEN-G_Paquete_de_revision_de_continuidad_R6.html
- **Regulation (EU) 2016/679 (GDPR)** https://eur-lex.europa.eu/eli/reg/2016/679/oj/eng
- **Area page** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/areas/operacion.html
- **Lite and Enterprise operation** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/52_SEVEN-G_Manual_de_operacion_de_IA.html#lite-and-enterprise-operation
- **Building solutions with AI** https://seachad.github.io/AI_CONSULTING/SEVEN-G/html/en/53_SEVEN-G_Construccion_de_soluciones_con_IA.html

LEGAL NOTICE

Legal notice: a generic, free methodological aid provided "as is" for information purposes; it is not legal advice and has not been legally reviewed for any case. Each organisation adapts it and is responsible for its legal review and its regulatory compliance; it may not reflect recent regulatory changes. To the fullest extent permitted by law, the author accepts no liability whatsoever for the effects of its application. It does not grant certification of any kind. Details in document 93.

CC BY 4.0 (content) · MIT (code) · Fernando García Varela · seachad.github.io/AI_CONSULTING/