

SEVEN-G

Company · Self-consulting

What SEVEN-G is, how the initiative funnel gets filled all the way to the board dashboard, and how to implement it in your company by your own means.

The problem it solves

- Many companies have AI pilots that work, but few reach production.
- Value gets declared ("we freed up 200 hours") and almost nobody checks it.
- Nobody has clear authority to stop an initiative that is not working.
- The board asks about AI and gets a list of technical projects, not business decisions.
- Is the company transforming with AI, or just a bit more efficient? Almost nobody can answer that with evidence.

This is not a technology problem: it is a governance and measurement problem.

What SEVEN-G is and what it delivers


- Seven-phase Enterprise Value & Governance: an enabling phase 0 and seven value phases, with governance throughout.
- It organises AI as a portfolio: every initiative has a phase, a cost, expected value and an owner.
- It places decision gates between phases: no verified evidence, no progress.
- It separates who decides, who builds and who controls, and turns regulation into phases, roles and evidence.
- It measures value in money and with a status —validated, declared or estimated— and shows it to the board on a dashboard.
- It is free and modifiable: documents, Word templates and tools, under an open licence.

Two levels: the company and each initiative

- C1 · Diagnosis: inventory, maturity across seven dimensions and the transformation index.
- C2 · Direction: the board approves the AI thesis, ambition per sphere and risk appetite.
- C3 · Portfolio: which initiatives get funded, and in what order.
- C4 · Oversight: board dashboard and the log of decisions and recommendations.
- C5 · Annual review: what worked, what did not and what changes.
- And, inside that portfolio, each initiative goes through its own eight-phase lifecycle.

An initiative's lifecycle

Eight phases with a decision gate at the end of each; below, the funnel stage where each phase shows up



What can happen at each gate

Five outcomes; whoever provides the evidence never verifies it or decides

Proceed	With conditions	Iterate	Pivot	Stop
All criteria are met and the evidence is verified: it moves to the next phase.	It proceeds, with a deadline and owner for what is pending. Never for critical security, legal or human-oversight controls.	The phase repeats with a specific adjustment and returns to the same gate. The third time, the higher body decides.	The hypothesis does not hold but another reasonable one exists: back to phase 2. Only at G1, G2 and G3.	No plausible value, or unacceptable risk. The reason and lessons are recorded. It is a good outcome.

At G7 the outcomes are Scale, Iterate or Retire.

■ FROM DOCUMENT TO DASHBOARD

How the funnel gets filled

Nobody draws the funnel or the dashboard: they come out of what is filled in and decided at each phase.
This is the path.

Phases 0 to 3: from idea to proof of concept


What is done, which templates are filled in and which gate closes each phase

F0 Context and constraints	F1 Discovery	F2 Value hypothesis	F3 Feasibility and risk
What is done. The idea is authorised: sponsor, owner, scope, constraints and intensity (Lite or Enterprise). Filled in. P01 Initiative charter · P02 Context and constraints · P03 Roles · P04 Intensity · P05 Inventory entry.	What is done. Opportunities are compared, including non-AI alternatives, and a sphere and ambition level are proposed. Filled in. P06 Opportunity portfolio and screening · P07 Sphere and ambition · P31 Use case sheet.	What is done. A falsifiable hypothesis in money, with baseline, attribution method and stop criteria set in advance. Filled in. P08 Value hypothesis canvas · P09 Baseline. Support: T11 calculator.	What is done. Proof of concept: feasibility, regulatory classification, risks and suppliers. It is the main stop gate. Filled in. P10 Feasibility · P11 Regulatory classification and impact · P12 Risks · P13 Mitigation · P14 Supplier.
G0 → enters "Proposed"	G1 → stays in "Proposed"	G2 → "Value hypothesis"	G3 → "Proof of concept"

Every gate decision is recorded in P29 and in the T01 register, with the status of each criterion and its evidence.

Phases 4 to 7: from design to production and retirement

Go-live does not close the cycle: it opens operation, with periodic reviews



Every gate decision is recorded in P29 and in the T01 register, with the status of each criterion and its evidence.

The initiative register (T01): where everything is recorded


Each initiative is entered like an opportunity in a CRM; here, the sample register with fictitious data



- One bar per phase: how many initiatives are in each one now and how much expected value they carry.
- When a gate is decided, the register stores the date, the outcome, who verified and who decided, and the status of each criterion.
- From those dates it works out time in phase by itself and flags as "stalled" those beyond the reference time limit.
- Figures carry a status: validated, declared or estimated. What is unknown stays "no data", never zero.
- Alerts: pending gates, overdue conditions, risks without an owner.
- Everything is saved in the company's own JSON file: no server and no data sent to anyone.

The portfolio funnel: how to read it

The AI portfolio, read as a sales funnel. Illustrative data



- Inside the funnel, only what is in progress: 5 proposed, 4 with a value hypothesis, 3 in proof of concept and 4 in development.
- On the left, how many reached each stage (40 → 31 → 24 → 17) and how long they take. The red triangle warns: proof of concept exceeds its time limit.
- On the right, those that did not pass each stage and why: no plausible value, hypothesis refuted, unacceptable risk, cost above value.
- Below, those that made it through: 11 in use, with their value and cost, and 1 disengaged.
- The goal is not to convert more: stopping an unviable initiative at G3 is a good outcome. 100% conversion would point to soft gates.

The board dashboard (T17): how to read it

It is generated from the same register, with nothing typed twice. Fictitious data



- 15 cases: how many are in use, in development, in proof of concept, proposed and out.
- Costs of €2.4M a year and total return of €7.3M (efficiencies plus return): annual net is €4.8M.
- The yellow card is the one that matters: only 36% of the value is validated by finance. The rest is declared or estimated.
- Released capacity that has not been materialised is shown separately and does not add up.
- In red: only 3 of 15 cases have all the controls (impact assessments, security, user manual).
- Other pages: funnel and lifecycle, history and adoption, risk and compliance, inventory.

If a figure on the dashboard looks wrong, the dashboard is not touched: the evidence is fixed in the register and it is generated again.

■ TWO EXAMPLES

One initiative that stops and one that reaches production

Fictitious, illustrative data: they do not describe any real company.

Example 1 · When an initiative stops

Candidate screening in recruitment (HR)

- Phases 0 and 1: HR sponsors it (P01–P05). Because it decides about people, intensity is Enterprise (P04). It enters "Proposed".
- Phase 2: the hypothesis (P08) promises 40% less screening time against the baseline (P09). The pilot confirms it: G2, Proceed.
- Phase 3: the regulatory classification (P11) makes it a high-risk system. The bias test finds it penalises older candidates.
- It iterates twice, adjusting variables: fixing the bias destroys the predictive power that justified the project.
- G3: high residual risk with no viable mitigation (P12, P13). The AI committee decides to Stop; it is recorded in P29 with its reason.
- In the funnel it shows to the right of "Proof of concept": Discarded · unacceptable risk. Lesson: test for bias earlier, in phase 2.

STOPPED AT G3 · BUILDING AND RUNNING AN UNVIABLE SYSTEM WAS AVOIDED

Initiative 1, on the lifecycle

It reached phase 3 and did not pass gate G3



STOPPED AT G3 · it leaves the funnel as "Discarded", with its reason

Example 2 · When an initiative reaches production

First-line IT incident response assistant

- Phases 0 and 1: IT sponsors it; against the non-AI alternative (more staff), the operations sphere with Optimise ambition is proposed (P07).
- Phase 2: hypothesis of 30% less resolution time for simple incidents; the controlled pilot yields 32%.
- Phase 3: no relevant bias, medium residual risk accepted by the right body (P12), supplier assessed (P14). G3, Proceed.
- Phases 4 and 5: mandatory human oversight above a severity level (P17), tested rollback (P19) and go-live sign-off (P23).
- Phase 6: fourteen months in use. Finance validates €1.1M a year (P28); released capacity was reassigned to complex incidents.
- In the funnel it sits under "In use"; on the dashboard it adds to annual net because its value is validated, not declared.

IN PRODUCTION · €1.1M OF VALIDATED ANNUAL RETURN · QUARTERLY R6

Initiative 2, on the lifecycle

It went through phases 0 to 6 and remains in use; phase 7 will come with G7



IN USE · each R6 confirms it keeps delivering value under control

How value is measured

- Every amount carries a formula and a status: validated (finance), declared (the team) or estimated.
- Released hours are not savings until they are materialised: a cost goes down or they are reassigned to a valuable activity.
- Cost is complete: build, consumption, operation, human oversight and retirement.
- The transformation index uses eight evidence signals to tell transformation apart from mere efficiency.
- Ten measurement rules (document 40). No figure enters the portfolio net without its status.

■ SELF-CONSULTING

Implementing it in your company, by your own means

The framework is designed to be applied without hiring anyone. This is what it takes, and in what order.

Before starting: what the company needs

- An implementation mandate signed by the top executive, with objective, perimeter and deadline (P32).
- The board knows the plan and reserves a session around week 12 to approve the direction (C2).
- An implementation lead with real dedication and a core team: business, technology, data, risk, compliance, security, people and finance.
- Internal audit informed and legal counsel available for regulatory classification.
- A register: T01, or a spreadsheet with the same fields.
- No moratorium: AI activity continues; only what the diagnosis flags as unacceptable risk is stopped.

Lite or Enterprise: proportionality

- Two different decisions: implementation scope belongs to the company; intensity, to each initiative.
- Size does not decide. One Enterprise criterion is enough for an initiative: high risk, decisions about people, direct exposure, acting agents, protected data, critical function, Transform or high investment.
- Minimum Lite path: board with a quarterly AI item, monthly executive committee and a one-person part-time AI office.
- In Lite, gates are grouped (G0–G2 together, G4–G5 together, six-monthly R6). G3 always stands alone.
- No mandatory evidence is skipped in Lite: it is simplified. The inventory is never reduced.

The 90-day plan

Three months, three results and one board decision

Month 1 · Diagnosis (C1)

Mandate (P32). Inventory of all AI, including third-party and unauthorised, with a completeness statement. Evidence-based maturity (T15) and transformation index (T14). Milestone 1: verified C1 report (P33).

Month 2 · Risks and opportunities

Workshops per sphere: risks of what already exists (P11, P12) and opportunities with estimated value (P06, P07). Draft AI thesis, ambition and risk appetite (P35). Milestone 2: C2 proposal ready for the board.

Month 3 · Governance and portfolio

Bodies and roles with no incompatibilities (P38, P03), gates and policies, first dashboard with real data (T17) and prioritised portfolio (P36). The board approves C2 and the committee, C3. Milestone 3: business as usual.

From week 13, no new initiative spends budget without going through G0.

Three guarantees that are never negotiable, even at small scale

- Whoever builds neither verifies nor decides on their own work.
- Risk sign-off comes from someone independent of the team.
- The board approves the direction and the risk appetite: this is not delegated.
- One person can hold several roles if they are not incompatible; the T01 register checks it (P03).
- What already exists is regularised with approved deadlines, not ignored: inventory first, then risks, then gates.

Common mistakes when implementing it yourself

- Applying Enterprise to everything: the framework becomes heavy and is abandoned. Or applying nothing "because we are small".
- Filling in templates after deciding: after-the-fact documentation invalidates the gate.
- Adding released hours as savings on the dashboard.
- Never stopping anything: a portfolio with no stops has formalities, not gates.
- An incomplete inventory: third-party AI and unauthorised use count too.
- Leaving the board out until the end: its session is booked on day one.

Where to start tomorrow

- One afternoon: open the sample register (T01) and its dashboard (T17), and answer the maturity diagnostic (T15) with your real situation.
- One week: documents 00, 01 and 90; decide Lite or Enterprise scope and draft the mandate (P32).
- Adoption can be modular: just the register, just the gates, just measurement or just the dashboard, and extend later.
- If you want support, you can ask any professional; document 91 explains how to avoid depending on them.
- None of this requires registering or hiring anyone.

Thank you

SEVEN-G is free: it is used, downloaded and adapted without registering, and the site will never write to you. If you want to talk to the author, the initiative is yours.

SEVEN-G is a registered trademark of Fernando García Varela · CC BY 4.0 (content) · MIT (code)

<https://www.linkedin.com/in/fernandogarciavarela/>