TEMPLATE P17 · GOVERNANCE AND HUMAN OVERSIGHT DESIGN

Governance and human oversight design

Establishes what the system decides, what a person validates and what is never delegated, with what autonomy level and with what human oversight modes.

Fernando García Varela Version 0.2 (working draft) 19-10-2026 Draft for review

IN THIS DOCUMENT

01 Instructions for use

02 Identification

03 Decision context

04 Autonomy level

05 Allocation of decisions

06 Human oversight modes

07 Effective capability of the overseer

08 Transparency and rights of affected persons

09 Responsibilities and escalation

10 Oversight indicators

11 Quality criteria

12 Approval and verification

13 Version control

Legal notice and disclaimer. SEVEN-G is a reference methodological framework provided "as is" and for information purposes only. It does not constitute legal, regulatory, financial or professional advice, nor does it guarantee compliance with any law or standard. References to general regulation (such as the EU AI Act, the GDPR, DORA or NIS2), to technical standards and to regulation specific to each sector or jurisdiction may be incomplete, may not apply to a particular case or may become out of date as a result of regulatory changes, interpretations or supervisory positions after their consultation date. **Each organisation that uses SEVEN-G is solely responsible for identifying the regulation that applies to it, verifying that it is current and certifying its own regulatory compliance**, with the appropriate qualified advice. This methodology is a generic, free aid shared with the community so that nobody has to start from scratch; each person or organisation can and should adapt it to its own use. It must not be inferred that its legally sensitive parts have been reviewed by legal counsel: those reviews, for each company or sector, are the ultimate responsibility of the company, consultant or organisation that uses it. Although every effort is made to keep it up to date, some regulation may have changed without being reflected here. To the fullest extent permitted by law, the author accepts no responsibility whatsoever for the effects of its application in any organisation or for its full applicability. The methodology does not grant certification of any kind. The author accepts no liability for any use made of this content or for any decisions taken on the basis of it. The data, figures, companies and cases in the examples are fictitious or illustrative.

01

Instructions for use

Aspect	Guidance
When it is used	Phase 4 (Solution design), before requesting G4. It is reviewed at G5 with what was observed in the pilot and at every R6, or earlier if the autonomy level or the affected population changes.
Who completes it	AI Product Owner and AI Technical Owner, with the business area that will use the system.
Who verifies it	AI Office under Lite; AI Auditor under Enterprise. The AI Risk Owner issues clearance.

Aspect	Guidance
At which *gate* it is reviewed	G4; its actual operation is checked at G5 (P22) and at R6.
Lite versus Enterprise	Lite may omit the fields marked **(Enterprise)**. A system at level A2 or A3 whose actions affect third parties, money, personal data or production systems is Enterprise (criterion "agents with the ability to act", 01 §9.2).
Relationship	P04 Intensity determination · P11 Regulatory classification · P18 Security design · P19 Rollback plan · P24 Operations manual · P25 Monitoring · document 35 (autonomy and AG controls) · document 52.

Rules:

- Human oversight controls do not allow **Proceed with conditions** at the *gate*: if they are not resolved, the result is Iterate (01 §7.3).
- Oversight is designed so that a person can **understand, correct, override and interrupt**, not merely so that they "are present".
- Any increase in the autonomy level or in the affected population should be treated as a relevant change (P27) and be subject to verification and a decision equivalent to G4.

This template does not constitute legal advice. Obligations regarding human oversight and automated decisions are confirmed with the regulatory mapping (document 34).

02

Identification

Field	Content	Guidance
Use case code		IA-AAAA-NNN (T01).
Use case name		
AI system(s)		Identifier in the inventory (T02).
Intensity		Lite · Enterprise.
Ambition level		Optimise · Augment · Transform.
Regulatory classification		Prohibited · High risk · Transparency obligations · Minimal risk · Out of scope · Pending classification.

Field	Content	Guidance
Document version		
Author		Role and name.
Date		DD-MM-AAAA.
Document status		Draft · Submitted for verification · Verified · Approved · Superseded.

03

Decision context

Field	Content	Guidance
Process or decision in which the system intervenes		Describe it in business language (P31).
Users		Who uses the output.
Affected persons		Who receives the effect: customers, employees, citizens, suppliers.
Does it significantly influence decisions about people?		Yes · No. If Yes, it is an Enterprise criterion.
Are there decisions based solely on automated processing with legal or similarly significant effects?		Yes · No · To be determined. If Yes or To be determined, data protection must analyse Article 22 of the GDPR.
Consequence of a system error		What happens if the output is incorrect and who bears it.
Reversibility of the decision or action		Reversible at no cost · Reversible at a cost · Irreversible.

04

Autonomy level

Tick a single level. The levels are defined in document 35.

Level	Name	What the system does	Human role	Tick
A0	Assistance	Informs, summarises or generates content.	The person decides and executes.	☐
A1	Recommendation	Proposes a specific decision or action.	The person validates each action before it is executed.	☐
A2	Supervised action	Executes actions within defined limits.	Supervises, can interrupt and reviews after the fact.	☐
A3	Autonomous action	Executes sequences of actions without individual review within strict limits.	Sets limits, supervises aggregates and has a kill switch.	☐

Field	Content	Guidance
Justification for the level chosen		Why a lower level is not enough. The lowest level that achieves the objective is preferable.
Level in the pilot		It may be lower than the target (for example, A1 in the pilot and A2 in production).
Conditions for moving up a level (Enterprise)		Evidence required (results, override rate, incidents) and the body that authorises it.
Conditions for moving down a level		Thresholds that require autonomy to be reduced; they are carried over to P19 as a degraded mode.

05

Allocation of decisions

One row per decision or action in the process. **Allocation:** The system decides · A person validates · Never delegated.

Decision or action	Allocation	Justification	Role that validates or decides	Information needed by the validator
Prioritise the queue of incoming requests *(illustrative example)*	The system decides	Reversible, with no direct effect on the person; processing remains human	—	—
Propose the resolution of a request *(illustrative example)*	A person validates	Affects the applicant; requires checking of documentation	Case handler	Proposed outcome, main reasons and documents consulted

Decision or action	Allocation	Justification	Role that validates or decides	Information needed by the validator
Reject a request ***(illustrative example)***	Never delegated	Significant effect on the person; requires human reasoning	Head of the area	Complete case file

Decisions that the company has declared non-delegable in its corporate policy (document 31) and that apply to this system:

Non-delegable decision	Reference in the policy	How it is technically guaranteed

06

Human oversight modes

Modes: Prior validation (a person approves each output before it takes effect) · Real-time oversight (a person observes and can intervene while the system acts) · Sample-based after-the-fact review · Aggregate oversight through indicators and thresholds.

Mode	Applies? (Yes · No)	Scope	Responsible role	Frequency or sample size	Intervention criterion	Where it is recorded
Prior validation						
Real-time oversight						
Sample-based after-the-fact review						
Aggregate oversight						

Guidance: A1 requires prior validation; A2 requires at least real-time oversight or after-the-fact review, and a kill switch; A3 requires aggregate oversight, strict limits (P18) and a tested kill switch (P19).

07

Effective capability of the overseer

Requirement	How it is guaranteed	Evidence
Understands the capabilities and limitations of the system		Training recorded in P20; limitations from P16.
Is aware of the risk of over-relying on the automated output		Training and interface design.
Can correctly interpret the output		Explanation or reasons shown alongside the output.
Can decide not to use, correct, override or reverse the output		Function available and logged.
Can interrupt the system		Kill switch (P18) and who can activate it.
Has time and a workload compatible with oversight (Enterprise)		Expected volume per person and time per review.

In high-risk systems, these requirements are checked against Article 14 of the EU AI Act (document 34).

08

Transparency and rights of affected persons

Field	Content	Guidance
Are people informed that they are interacting with an AI system or that the content is generated?		Yes · No · Not applicable. How and at what point.
Information on the logic and the reasons for the output		What is explained and to whom.
Channel for requesting human review or contesting		Channel, internal response time limit and role that reviews.
Log of complaints related to the system (Enterprise)		Where they are recorded and how they feed P25 and P27.

Field	Content	Guidance
Immediate handover to a person (direct exposure)		Signs of vulnerability that trigger it, express request, channel and maximum time until a person responds (35 §4.6).
Log of handovers		Where reason, time and time to response are recorded; they are carried over to P25.

Notice texts and register in P49.

09

Responsibilities and escalation

Role	Holder	Deputy	Oversight responsibility	Escalates to
AI Product Owner			Design of the allocation of decisions and of adoption	AI Sponsor
AI Operations Owner			Day-to-day operation of oversight and alerts	AI Committee in S1–S2 incidents (P26)
Designated overseers			Validation, review and override	AI Operations Owner
AI Risk Owner			Clearance on the sufficiency of oversight	AI Committee

10

Oversight indicators

They are carried over to P25 with their thresholds.

Indicator	Definition	Alert threshold	Frequency
Human override rate	Outputs corrected or rejected ÷ outputs reviewed		
Average validation time	Time between proposal and validation		

Indicator	Definition	Alert threshold	Frequency
After-the-fact review coverage	Outputs reviewed ÷ outputs produced		
Acceptance without effective review (Enterprise)	Validations taking less than the reasonable minimum time		
Complaints from affected persons	Number per period		
Handovers to a person	Handovers by reason and requests not served		

11

Quality criteria

The formal G4 criteria on human oversight are in document 21 (and its checklist LV-G4 in document 22); autonomy controls are in the AG controls of document 35.

#	Check	Status (Met · Not met · Not applicable · Pending)
1	A single autonomy level is assigned and justified.	
2	Each relevant decision in the process has an allocation, and the non-delegable ones are identified.	
3	The oversight modes are consistent with the autonomy level.	
4	The overseer can override and interrupt, and is trained to do so.	
5	There is a human review channel for affected persons where appropriate.	
6	The oversight indicators have a threshold and have been carried over to P25.	
7	Data protection has reviewed the applicability of Article 22 of the GDPR where appropriate.	

12

Approval and verification

Function	Role	Name	Date	Signature or reference
Author	AI Product Owner			
Co-author	AI Technical Owner			
Clearance	AI Risk Owner			
Verifier	AI Office (Lite) · AI Auditor (Enterprise)			
***Gate* decision-maker**	According to 01 §7.5 (recorded in P29)			

Segregation of duties: the AI Risk Owner and the verifier are not part of the team that builds; no one verifies or approves their own work.

13

Version control

Version	Date	Changes
0.1	16-09-2026	First version of the template.
0.2	19-10-2026	Immediate handover to a person and handover log (8) and handover indicator (10) (35 §4.6) (D159).

External sources cited

Links to the official or primary source, checked before publication. Always verify the version in force on the date of use.

Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation)

· https://eur-lex.europa.eu/eli/reg/2016/679/oj/eng · Parlamento Europeo y Consejo de la UE · 27-04-2016

Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act) · https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng · Parlamento Europeo y Consejo de la UE · 13-06-2024