

TEMPLATE P18 · SECURITY DESIGN (INCLUDING AGENTS)

Security design (including agents)

Documents the threats to the AI system and the designed controls, with particular attention to the identity, permissions and limits of agents that act.

IN THIS DOCUMENT

01	Instructions for use	08	Human validation of sensitive actions
02	Identification	09	Kill switch
03	Threats and controls	10	Action log
04	Agent identity	11	Security testing, including prompt injection
05	Least permissions	12	Quality criteria
06	Intent-based access control	13	Approval and verification
07	Action limits	14	Version control

Legal notice and disclaimer. SEVEN-G is a reference methodological framework provided "as is" and for information purposes only. It does not constitute legal, regulatory, financial or professional advice, nor does it guarantee compliance with any law or standard. References to general regulation (such as the EU AI Act, the GDPR, DORA or NIS2), to technical standards and to regulation specific to each sector or jurisdiction may be incomplete, may not apply to a particular case or may become out of date as a result of regulatory changes, interpretations or supervisory positions after their consultation date. **Each organisation that uses SEVEN-G is solely responsible for identifying the regulation that applies to it, verifying that it is current and certifying its own regulatory compliance,** with the appropriate qualified advice. The author accepts no liability for any use made of this content or for any decisions taken on the basis of it. The data, figures, companies and cases in the examples are fictitious or illustrative.

01

Instructions for use

Aspect	Guidance
When it is used	Phase 4 (control design). The tests in block 11 are planned in phase 4 and executed in phase 5, before G5. It is reviewed at R6 and upon changes to tools, permissions or model.
Who completes it	AI Technical Owner, with information security.
Who verifies it	AI Auditor under Enterprise; AI Office under Lite, with a review by information security.
At which <i>gate</i> it is reviewed	G4 (design) and G5 (tests passed and security sign-off in P23).

Aspect	Guidance
Lite versus Enterprise	Lite may omit the fields marked (Enterprise) . If the system does not execute actions (A0 or A1 according to P17), blocks 5 to 9 may be marked <i>Not applicable</i> with a justification; technical identity (block 4) and logging (block 10) are always completed.
Tool	T10 Agent security assessment.
Relationship	P15 Architecture · P16 Lineage · P17 Human oversight · P19 Rollback · P25 Monitoring · P26 Incident response · document 35 (AG controls) · document 36 (third parties).

Rules:

- Controls are referenced using the **AG codes of document 35**. No local codes are created in that format; the company's own controls are referenced with their internal code.
- Critical security controls do not allow **Proceed with conditions** at the *gate* (01 §7.3).
- The content the system reads (documents, emails, web pages, tool responses) is **data, not instructions**: the design must prevent that content from ordering actions.

02

Identification

Field	Content	Guidance
Initiative code	IA-AAAA-NNN (T01).	
AI system(s)	Identifier in the inventory (T02).	
Technology	Predictive ML · Generative AI · Agent · Language and document processing · Vision · Optimisation · Embedded third-party AI · Rules (not AI).	
Exposure	Internal · Employees · Customers indirectly · Customers or external persons directly.	
Autonomy level	A0 · A1 · A2 · A3 (P17).	
Intensity	Lite · Enterprise.	
Document version		
Author	Role and name.	
Date	DD-MM-AAAA.	

Field	Content	Guidance
Document status	Draft · Submitted for verification · Verified · Approved · Superseded.	

03

Threats and controls

Control status: Designed · Implemented · Tested · Pending · Not applicable.

Threat	Applies? (Yes · No)	Planned control	Reference (AG control of document 35 or internal control)	Owner	Status
Direct prompt injection					
Indirect prompt injection (through documents, emails, web or tools)					
Information leakage in responses or logs					
Poisoning of training data or the knowledge base					
Model extraction or manipulation					
Inputs designed to deceive the model					
Excessive permissions of the system or agent					
Unauthorised actions or actions outside the purpose					
Consumption and cost abuse					
Supply chain (models, libraries, suppliers)					
Lack of traceability of actions and intents					

Threat	Applies? (Yes · No)	Planned control	Reference (AG control of document 35 or internal control)	Owner	Status
Impersonation or synthetic content used against the process (Enterprise)					

04

Agent identity

Field	Content	Guidance
Own technical identity		Each agent has its own identity, not shared with people or other agents.
Accountable owner		Person who is accountable for the identity (normally the AI Operations Owner).
On whose behalf it acts		On behalf of a specific user (delegation) · On behalf of the company (process). If it acts on behalf of a user, never with more permissions than that user.
Credential management		Secrets vault; no credentials in prompts or code; rotation.
Revocation		How and how quickly credentials are revoked; who can do so.
Periodic review of the identity and its permissions (Enterprise)		Frequency and owner. Recorded in the non-human identity inventory (P54).

05

Least permissions

Operation: Read · Write · Delete · External sending · Financial transaction · Configuration change.

Tool or system	Operation	Specific scope	Justification	Prior human validation (Yes · No)
Order management system (<i>illustrative example</i>)	Read	Orders of the customer who starts the conversation	Check status	No

Tool or system	Operation	Specific scope	Justification	Prior human validation (Yes · No)
Order management system (<i>illustrative example</i>)	Write	Change of delivery address before dispatch	Resolve a frequent request	No, within the limits of block 7
Returns gateway (<i>illustrative example</i>)	Financial transaction	Refunds for the customer's own orders	Resolve returns	Yes, above the limit

Any permission not listed is denied.

06

Intent-based access control

Field	Content	Guidance
Authorised purposes of the agent		Closed list of tasks for which it may act.
Recording of the intent per action		Each action records the task or request that originates it.
Consistency check between intent and action		How an action that does not correspond to the declared purpose is blocked (for example, an external sending during an enquiry).
Handling of instructions contained in data		How legitimate instructions and content read are kept separate.
Chained actions (Enterprise)		Limit on steps and a check before each action with an external effect.

07

Action limits

Limit	Value	What happens when it is reached (Block · Human validation · Alert · Stop)
Maximum amount per action		

Limit	Value	What happens when it is reached (Block · Human validation · Alert · Stop)
Cumulative amount per period		
Number of actions per period		
Permitted recipients or domains		
Prohibited systems and data		
Maximum consumption cost per period		
Maximum number of steps per task		
Time window for action (Enterprise)		

08

Human validation of sensitive actions

Sensitive action	Sensitivity criterion	Who validates	Information seen by the validator	Maximum response time	If there is no response
Refund above the limit (<i>illustrative example</i>)	Amount	Customer service supervisor	Order, reason, history and the agent's proposal	4 working hours	It is cancelled and the customer is informed

By default, a sensitive action without validation **is cancelled**; it is never executed because the time limit has expired. It must be consistent with the allocation of decisions in P17.

09

Kill switch

Field	Content	Guidance
Mechanism	How the system is stopped (function, configuration, revocation of credentials).	

Field	Content	Guidance
Who can activate it		Roles, including at least one available outside working hours if the system operates continuously.
What it stops		New actions, actions in progress, pending queues.
Target time to effective stop		Measured in the test in block 11.
Safe state after the stop		Manual process or degraded mode (P19).
Authorisation to reactivate		A role different from the one operating the system when it was activated due to an incident.

10

Action log

Data logged	Is it logged? (Yes · No)	Observations
Date and time		
Agent identity and version of model and prompts		
User or process on whose behalf it acts		
Originating intent or task		
Tool, operation and parameters		Without exposing unnecessary sensitive data.
Result of the action		
Human validation (who, when, outcome)		
Blocks due to limits or intent-based access control		

Field	Content	Guidance
Retention period		Not shorter than that required by applicable regulation (document 34).

Field	Content	Guidance
Protection against tampering (Enterprise)		How the integrity of the log is guaranteed.
Access to the log		Who can consult it; the agent cannot modify it.

11

Security testing, including prompt injection

Type: Direct injection · Indirect injection · Information leakage · Privilege escalation · Limit evasion · Consumption abuse · Kill switch test · Other. **Result:** Passed · Passed with observations · Failed · Pending.

Test	Type	Date	Performed by	Result	Findings	Action and owner
Document with hidden instructions asking for data to be forwarded (<i>illustrative example</i>)	Indirect injection		Security team external to the project	Passed with observations	The agent did not forward the data, but quoted the instruction in its response	Filter quoted content · AI Technical Owner

Tests should be performed by people outside the team that builds. A *Failed* test on a critical control prevents the security sign-off in P23. The plan and report of adversarial testing campaigns are documented in P53.

12

Quality criteria

Formal G4 and G5 criteria in document 21; agent controls in the AG controls of document 35.

#	Check	Status (Met · Not met · Not applicable · Pending)
1	All applicable threats have a control and an owner.	
2	The agent has its own revocable identity.	

#	Check
3	Permissions are limited to those listed and are justified.
4	There is intent-based access control and there are limits with a defined consequence.
5	Sensitive actions require human validation and are cancelled without it.
6	The kill switch has been tested and the measured time meets the target.
7	The action log makes it possible to reconstruct who did what, why and with what result.
8	Injection tests have been performed before G5 and critical findings have been resolved.

13

Approval and verification

Function	Role	Name	Date	Signature or reference
Author	AI Technical Owner			
Review	Information security			
Clearance	AI Risk Owner			
Verifier	AI Office (Lite) · AI Auditor (Enterprise)			
Gate decision-maker	According to 01 §7.5 (recorded in P29)			

Segregation of duties: information security and the verifier are not part of the team that builds.

Version control

Version	Date	Changes
0.1	16-09-2026	First version of the template.